



PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013
Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064



AI-ENHANCED CAREER GUIDANCE SYSTEM FOR PERSONALIZED CAREER PATHWAYS

A PROJECT REPORT

Submitted by

ADITHYAN S NAMBIAR- 20221CSE0223

S KRISHNA KUMAR- 20221CSE0184

K NANDAN- 20221CSE0179

Under the guidance of,

Dr. RIYAZULLA RAHAMAN J

BACHELOR OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING

PRESIDENCY UNIVERSITY

BENGALURU

DECEMBER 2025



PRESIDENCY UNIVERSITY

Private University Estd. In Karnataka State by Act No. 41 of 2013
Itgalpura, Rajankunte, Yelahanka, Bengaluru - 560064



PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

BONAFIDE CERTIFICATE

Certified that this report AI-ENHANCED CAREER GUIDANCE SYSTEM FOR PERSONALIZED CAREER PATHWAYS is a bonafide work of ADITHYAN S NAMBIAR (20221CSE0223), S KRISHNA KUMAR (20221CSE0184), K NANDAN (20221CSE0179), who have successfully carried out the project work and submitted the report for partial fulfilment of the requirements for the award of the degree of BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE ENGINEERING, during 2025-26.

Dr. Riyazulla Rahaman J
Project Guide
PSCS
Presidency University

Mr. Muthuraju V
Program Project
Coordinator
PSCS
Presidency University

Dr. Sampath A K
Dr. Geetha A
School Project
Coordinators
PSCS
Presidency

Dr. Blessed Prince P
Head of the Department
PSCS
Presidency University

Dr. Duraipandian
Dean
PSCS & PSIS
Presidency University

Dr. Shakkeera L
Associate Dean
PSCS
Presidency University

Name and Signature of the Examiners

1) Pushpalatha M

Pushpalatha 21/12/25

2)

J. Saray 03.12.2025

PRESIDENCY UNIVERSITY

PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

DECLARATION

We the students of final year B.Tech in COMPUTER SCIENCE ENGINEERING, at Presidency University, Bengaluru, named ADITHYAN S NAMBIAR, S KRISHNA KUMAR, K NANDAN hereby declare that the project work titled **AI-ENHANCED CAREER GUIDANCE SYSTEM FOR PERSONALIZED CAREER PATHWAYS** has been independently carried out by us and submitted in partial fulfilment for the award of the degree of B.Tech in COMPUTER SCIENCE ENGINEERING, during the academic year of 2025-26. Further, the matter embodied in the project has not been submitted previously by anybody for the award of any Degree or Diploma to any other institution.

ADITHYAN S NAMBIAR USN: 20221CSE0223

S KRISHNA KUMAR USN: 20221CSE0184

K NANDAN USN: 20221CSE0179

PLACE: BENGALURU

DATE: 02-December 2025

ACKNOWLEDGEMENT

For completing this project work, We/I have received the support and the guidance from many people whom I would like to mention with deep sense of gratitude and indebtedness. We extend our gratitude to our beloved **Chancellor, Pro-Vice Chancellor, and Registrar** for their support and encouragement in completion of the project.

I would like to sincerely thank my internal guide **Dr. Riyazulla Rahaman J , Senior Scale Assistant Professor**, Presidency School of Computer Science and Engineering, Presidency University, for his/her moral support, motivation, timely guidance and encouragement provided to us during the period of our project work.

I am also thankful to **Dr. Blessed Prince, Professor, Head of the Department, Presidency School of Computer Science and Engineering Presidency University**, for his mentorship and encouragement.

We express our cordial thanks to **Dr. Duraipandian N**, Dean PSCS & PSIS, **Dr. Shakkeera L**, Associate Dean, Presidency School of computer Science and Engineering and the Management of Presidency University for providing the required facilities and intellectually stimulating environment that aided in the completion of my project work.

We are grateful to **Dr. Sampath A K, and Dr. Geetha A, PSCS** Project Coordinators, **Dr. Sharmasti vali, Program Project Coordinator**, Presidency School of Computer Science and Engineering, or facilitating problem statements, coordinating reviews, monitoring progress, and providing their valuable support and guidance.

We are also grateful to Teaching and Non-Teaching staff of Presidency School of Computer Science and Engineering and also staff from other departments who have extended their valuable help and cooperation.

ADITHYAN S NAMBIAR
S KRISHNA KUMAR
K NANDAN

Abstract

Selecting a suitable career path has grown more difficult for students and working professionals in the quickly changing professional landscape of today. Making decisions can be overwhelming due to the variety of career options, the rise of new roles, and the ever-changing job market trends. Even though they are helpful, traditional career counseling techniques are generally generic, time-consuming, and unable to provide the level of customization and flexibility required for effective guidance. Additionally, they do not incorporate real-time labor market insights, which could result in recommendations that fall short of industry standards or opportunities in the future.

In order to provide tailored, data-driven recommendations, this paper suggests an AI-Enhanced Career Guidance System that integrates machine learning and natural language processing (NLP). Academic performance, interests, and skills are among the user-specific data that the system gathers and compares to job descriptions, industry standards, and new skill trends. While machine learning algorithms analyze patterns and produce career pathways that match employability standards and individual goals, natural language processing (NLP) is used to process user inputs and contextual factors.

Comparing the system to traditional approaches, experimental evaluation shows that it greatly increases the accuracy and applicability of career recommendations. Users expressed greater satisfaction, pointing out that the system increased their confidence in making decisions and provided useful insights. The system facilitates effective career planning and closes the gap between academic preparation and practical opportunities by guaranteeing flexibility and scalability.

To provide more comprehensive career guidance in the future, the framework can be improved with psychometric tests, analytics of the global labor market, and sophisticated recommendation models. These AI-powered solutions have the potential to enable professionals, students, and organizations to make well-informed, sustainable career decisions that meet changing labor market demands.

Keywords— AI, Machine Learning, Career Guidance, NLP, Career Pathways

Table of Content

Sl. No.	Title	Page No.
	Declaration	
	Acknowledgement	I
	Abstract	II
	List of Figures	V-VI
	List of Tables	VI
	Abbreviations	VII
1.	Introduction 1.1 Background 1.2 Statistics of project 1.3 Prior existing technologies 1.4 Proposed approach 1.5 Objectives 1.6 SDGs 1.7 Overview of project report	1-9
2.	Literature review	10-13
3.	Methodology	14-26
4.	Project management 4.1 Project timeline 4.2 Risk analysis 4.3 Project budget	27-30
5.	Analysis and Design 5.1 Requirements 5.2 Block Diagram 5.3 System Flow Chart 5.4 Choosing devices 5.5 Designing units 5.6 Standards 5.7 Mapping with IoTWF reference model layers 5.8 Domain model specification	31-52

	5.9 Communication model 5.10 IoT deployment level 5.11 Functional view 5.12 Mapping IoT deployment level with functional view 5.13 Operational view 5.14 Other Design	
6.	Hardware, Software and Simulation 6.1 Hardware 6.2 Software development tools 6.3 Software code 6.4 Simulation	53-59
7.	Evaluation and Results 7.1 Test points 7.2 Test plan 7.3 Test result 7.4 Insights	60-68
8.	Social, Legal, Ethical, Sustainability and Safety Aspects 8.1 Social aspects 8.2 Legal aspects 8.3 Ethical aspects 8.4 Sustainability aspects 8.5 Safety aspects	69-73
9.	Conclusion	74-75
	References	76
	Base Paper	77
	Appendix	78-83

List of Figures

FIGURE	CAPTION	PAGE NO.
Fig 1.1	Sustainable Development Goals	6
Fig 3.1	The V model methodology	20
Fig 3.2	The W model methodology	21
Fig 3.3	The Devops methodology	21
Fig 3.4	Onion methodology	22
Fig 3.5	SDLC phases	23
Fig 3.6	Summary of Waterfall methodology	24
Fig 3.7	Agile/Scrum methodologies	25
Fig 3.8	Summary of project breakdown to task	26
Fig 4.1	project timeline	27
Fig 4.2	PESTEL analysis	28
Fig 4.3	Project phase risk matrix	29
Fig 4.4	project budget	30
Fig 5.1	Functional block diagram	33
Fig 5.2	System flow chart	34
Fig 5.3	interfacing a one unit	39
Fig 5.4	computing the values	41
Fig 5.5	The IoT World Forum Reference Model	43
Fig 5.6	Domain Model	45
Fig 5.7	Communication model	46
Fig 5.8	IoT deployment level	47
Fig 5.9	Functional view	48
Fig 5.10	Mapping IoT deployment level with functional view	50
Fig 5.11	Operational view	51
Fig 5.12	Process specification	52
Fig 7.1	API Connectivity Check	60
Fig 7.2	System Latency vs. Number of Users	64
Fig 7.3	Accuracy Distribution of Career Recommendations	64
Fig 7.4	Database Query Time Histogram	65
Fig 7.5	User Interaction Flow Success Rate	65

Fig B1	Meteor Springer Paper Acceptance email	78
Fig B2	Turnitin Similarity Report	79
Fig D1	Fig Landing (1)	80
Fig D2	Fig Landing (2)	81
Fig D3	Fig Login	81
Fig D4	Fig SignUp	82
Fig D5	Fig Real-Time Dashboard (1)	82
Fig D6	Fig Real-Time Dashboard (2)	83

List of Tables

TABLES	CAPTION	PAGE NO.
Table 2.1	Summary of Literature reviews	11-12
Table 5.1	Summarizing requirements	32
Table 5.2	Choosing devices	35-36
Table 5.3	Choosing Processor	36-37
Table 5.4	SBERT Model Table	37-38
Table 5.5	Data Layer	38
Table 5.6	Application Layer	39
Table 5.7	Mapping Project layers with IoTWFRM	44
Table 5.8	Domain Model Specification	45
Table 6.1	Additional Supporting Tools	56-57
Table 7.1	Test points Table	60
Table 7.2	Test Plan Table	61
Table 7.3	Test Methodologies Table	61-62
Table 7.4	Test Result Table	62-63
Table 7.5	Performance Characteristics Table	63
Table 7.6	Reasons for Performance Variations	67
Table 7.7	Evaluation of Functional Units	67
Table 8.1	Key Sustainable Design Principles	72

Abbreviations

• AI — <i>Artificial Intelligence</i>
• CGS — <i>Career Guidance System</i>
• LMI — <i>Labour Market Intelligence</i>
• ML — <i>Machine Learning</i>
• NLP — <i>Natural Language Processing</i>
• SBERT — <i>Sentence Bidirectional Encoder Representations from Transformers</i>

Chapter 1

INTRODUCTION

The process of finding the most suitable career path is now increasingly difficult as a result of rapidly changing technologies, increased globalization and increasing demands from employers for new skills and abilities. Traditionally, many forms of career counseling were based upon static surveys, assessments of an individual's personality type and/or a "keyword" search for jobs based on the input provided by the user. The problem with traditional approaches is they are unable to fully understand the deeper meaning of user interest and the true technical and other skill needs of many of today's work environments, resulting in generic and incorrect suggestions. Today, students/learners need a more personalized and dynamic form of career counseling that takes into account both their own unique characteristics/profiles and the current state of the labor market.

Our project will provide the solution to the issues listed above using the AI-Enhanced Career Counseling System utilizing SBERT technology for analyzing the semantic meaning of user data rather than simply matching keywords. SBERT allows for the analysis of the semantic meaning of job postings and the semantic meaning of a user profile, therefore enabling the intelligent matching of user profiles to job postings. This is accomplished by generating context-based representations (embeddings) of the job posting and the user profile and then comparing them to calculate the similarities. Our system utilizes a Flask-based back-end and MongoDB Atlas for scalable and cloud-based data storage to create a semantic recommendation engine to enable the user to be presented with high-relevance career options based on their unique characteristics/profiles. The objective of this study is to modernize the practice of career counseling through the use of deep learning-driven insights to ensure that users receive career counseling that aligns with the current needs of today's industries.

1.1 Background

In the current fast-paced world, choosing a fitting career has turned out to be more and more complicated. The vast amount of newly created job positions along with the continuous advancement in technology and the changing nature of the labor market, are making it really hard for students and professionals to pinpoint the way that perfectly fits their interests, skills, and aspirations for the future. The lack of clarity along with the huge number of options

available is causing the already difficult task of making a career choice to become even more so. The traditional methods for career counselling make use of unchanging questionnaires, general personality tests, and standard evaluations. Although they are helpful, these methods take a lot of time and usually do not lead to the delivery of real-time market trend-based and personalized advice. Consequently, many of the people involved are provided with wide and outdated advice that does not correspond to the current industries' dynamic needs. In order to eliminate these constraints, modern techniques are being adopted that are increasingly using Artificial Intelligence (AI) technologies such as Natural Language Processing (NLP) and Machine Learning (ML). These methods make it possible to create adaptive, data-driven systems that can not only analyze user profiles but also understand the meaning semantically and connect personal strengths with suitable career opportunities in an intelligent way. Thus, career counseling can become more precise, quicker, and more effective by incorporating AI-based insights—leading to better decision-making and more successful career results.

1.2 Statistics

India faces a severe challenge in employability despite its ability to produce a large number of graduates every year. National reports and studies reveal that only about 42.6% of Indian graduates are regarded as employable in both technical and non-technical job roles. This scenario indicates a huge gap in skills between the classroom and the industry. On the other hand, the gap calls for modern and targeted career guidance approaches that not only train students in the right skills but also connect education to the professional needs. Regional employment trends have also shown major inequalities. The state of Karnataka has about the same employability percentage as the national average, which indicates that the problem still exists even in a scenario with good educational and industrial synergies. Bengaluru, on the other hand, still has its status as one of the main cities providing jobs in India thanks to its high availability of job openings in such areas as IT, AI/ML, finance, biotechnology, and product engineering. The problem that comes with the influx of graduates into the market is that it leads to a situation where they have to compete with many others, which is very challenging for an individual to get his or her place. Therefore, it is imperative to have professional and large-scale AI-driven career counseling systems that will help students to recognize their strengths and align them with the labor market's rapidly changing demands.

1.3 Prior Existing Technologies

Artificial intelligence technology has given rise to numerous initiatives that aim to enhance career planning and talent mapping. An example of this is the Salesforce company, which has created AI-based solutions like Career Connect and Career Agent that scrutinize employee profiles and suggest career paths, training modules, and internal opportunities tailored to an individual. Such platforms take advantage of sophisticated AI models to determine strengths and weaknesses in terms of skills, hence facilitating the movement of employees within an organization. Research indicates that this kind of technology has a substantial contribution to employee engagement and internal talent development thereby reflecting the good influence of automated career guidance systems. Moreover, machine learning-based career recommendation systems that couple users to appropriate job roles have been researched significantly in academia. The latter are selected on the basis of a user's skills, educational background, and interests. Among the open-source initiatives, one can find the Machine Learning driven recommendation engines on GitHub which apply classification and clustering algorithms to forecast optimal career paths. These systems are the mainstay for the research of intelligent guidance mechanisms, but they are hindered by common limitations in comprehending the context due to their dependence on keyword or statistical models. This limitation provides the rationale for the application of more sophisticated semantic technologies—like SBERT—so as to enhance the precision and individualization of career recommendations.

1.4 Proposed Approach

Project Concept: The project has its main objective, which is to create an AI-Enhanced Career Guidance System that will give personalized and data-driven career recommendations. The system will analyze the individual's academic background, skills, and interests, and will be continuously updated with the labour market trends. The aim is to help the users select the career paths that are most suitable for them with more confidence and clarity.

Rationale: Choosing a career is a difficult task for students and working people primarily because of the enormous number of opportunities that are constantly being created and the equally fast-changing expectations of different industries. The methods of counselling that are in use today are no match for the constant changes of the times; besides, they are not able to

provide customized insights either. Thus, the need for an artificial intelligent system that can do the guiding and would-be users that can be easily persuaded to make informed career decisions is very much felt.

Methodology: The system relies on Natural Language Processing (NLP) for the interpretation of qualitative user inputs, which are user's preferences, interests, and experience descriptions, and on Machine Learning (ML) for user profile analysis and relevant career pathway generation. The output of the model is based on the fusion of semantic understanding, skill mapping, and labour market intelligence through which the recommendations are made tailored to one's personal goals along with the industry's demand and employability requirements.

Expected Use: Career guidance that is personalized according to one's needs for students as well as working professionals. Educational institutions provided with support in academic and placement counselling. HR departments and recruitment platforms receiving assistance in employee skill development. Career transitions and upskilling are parts of the process whereby people are assured to make informed decisions. **Disadvantages of the Method:** The main factor that determines the accuracy is the extent and quality of the data provided by the user. Might not be able to adapt immediately to the sudden fluctuations in the market.

1.5 Objectives

Objective 1: To create an engine for intelligent career profiling and semantic understanding Engine

Create a dialogue system that can acquire user characteristics like academic background, technical and soft skills, domain interests, and personal preferences with a target of $>95\%$ correct intent recognition through NLP parsing for the whole interaction. Use SBERT semantic embedding models (all-MiniLM-L6-v2) to represent both the inputs from users and the job descriptions in 384-dimensional vectors allowing the application of cosine similarity for similarity measurement with a response time of less than 200 ms for each query, thus providing the real-time generation of recommendations.

Objective 2: To put in place a pipeline for career recommendation based on ML

Train and adjust SBERT with a career-domain corpus aiming at the $\text{NDCG}@10 \geq 0.80$, $\text{MRR} \geq 0.85$, and $\text{Precision}@5 \geq 0.75$ as the recommendation relevance indicators. Build a ranking engine based on the Top-K nearest-neighbor search over vector embeddings and conduct

testing using 80/20 train-test split along with k-fold cross-validation (k=5) to get a performance that is generalizable. Incorporate data on labor market trends collected weekly to make sure that the recommendations are in line with the new skills and job roles.

Objective 3: Provide a secure and scalable system and user data management

Utilize Flask microservices and MongoDB Atlas cloud storage to deploy the system allowing the processing of over 10,000 user requests at the same time and having the 95th percentile response time not exceeding 300 ms. Provide role-based access control (RBAC) for students, counselors, and administrators. Employ AES-256 encryption for personal sensitive data and TLS 1.3 for secure communication, being compliant with data protection policies, and ensuring an uptime of the system of at least 99.5% thanks to the automatic data synchronization methods.

Objective 4: Build User-Centric UI and RAG-Enabled Conversational Assistant

To start with, develop a web UI that is not only interactive but also enables real-time result visualization, recommended learning paths, job market dashboards, and skill-gap analytics with a latency of UI response of less than 150 ms for each interaction. The incorporation of a Retrieval-Augmented Generation (RAG) module is to acquire data that is career-specific and to produce personalized recommendations which are based on reasoning. Users will have the option to select from various output formats like downloadable reports in PDF or CSV formats and also have access to progress-tracking dashboards.

Objective 5: Demonstrate Practical Impact, Accessibility, and Deployment Feasibility

The system will be deployed as a web-based application that will be accessible to students, professionals, and academic counselors across various platforms with the capability of over 1,000 recommendations daily and server resource usage of less than 50% of CPU and 65% of RAM on a modest infrastructure of 2 vCPU and 4GB RAM. Along with the deployment of user satisfaction verification through a pilot test that will aim at getting a positive feedback rate of 85% or more based on the criteria of recommendation relevance, novelty, and ease of use. The delivery will consist of deployment and maintenance documentation that will make the institution adoption feasible within 1-2 hours installation time without the need of specialized expertise.

1.6 SDGs

The **AI-Enhanced Career Guidance System** aligns with key UN Sustainable Development Goals by improving access to quality education and supporting decent work opportunities. It helps reduce skill gaps and promotes equal career opportunities. *Figure 1.1 shows the 17 UN SDGs adopted in 2015.*



Fig. 1.1 Sustainable Development Goals

SDG 4: Quality Education

The Quality Education objective of the United Nations is directly correlated to the AI-Enhanced Career Guidance System by the fact that it allows personalized learning and career support to be distributed equally. The system uses AI-based semantic analysis and data-driven insights to guide students towards the skills they need to develop and provides structured academic enabling pathways that are in sync with their growth. Its real-time guidance and reduced reliance on the limited number of human counsellors mean it can nicely fit into any social or geographical background and still offer quality education that is inclusive. The platform helps decision making of the learners and eventually strengthens the academic as well as the employability outcomes.

SDG 8: Decent Work and Economic Growth

The system is a strong supporter of SDG 8 by enhancing employment readiness and AI-supported career recommendations based on actual labour market trends. Through the matching of personal strengths with relevant job opportunities and emerging skill requirements, the system effectively reduces graduates' unemployment and underemployment. The platform provides a way to employability by directing users towards the learning pathways, certifications and opportunities for growth that are aligned with the industry. The real-time market awareness facilitates the education and industry gap closing, thus contributing to economic productivity and sustainable development of careers.

SDG 10: Reduced Inequalities

The project aligns with SDG 10 by giving equal opportunities to students of various socio-economic backgrounds to receive professional career guidance. Through the use of an AI-based counselling that is scalable and not only expensive private services, the platform makes sure that excellent guidance is provided for everyone, even the regions and institutions with limited counselling facilities that are underserved. This lessens the gap in employment opportunities, making it fairer to develop talents, and thus social inclusion is supported.

SDG 9: Industry, Innovation and Infrastructure

The system is a supporter of SDG 9 because of the application of such advanced technologies as NLP, Machine Learning, and SBERT-based semantic matching that are used to bring about the modernization of the old career guidance frameworks. The large user population can be supported through the cloud-based architecture that is scalable and that also enables the portals for institutions and recruitment platforms integration. By facilitating digital innovation and data-driven decision making, the system is reinforcing the technological infrastructure that is aligned with the long-term educational and workforce transformation.

1.7 Overview of project report

The Chapter 1 serves as an introduction, providing the context, background, and motivation for the development of the AI-powered, SBERT-enhanced career guidance system. It demonstrates the extent of the difficulties encountered by students and graduates by presenting employability statistics, surveys the technologies used for career counseling before and their traditional

drawbacks, introduces the semantic recommendation strategy, sets the project goals, and aligns the contributions with the UN Sustainable Development Goals.

The Chapter 2 comprises an extensive literature survey, which studies the academic literature and commercial systems in AI-based career guidance, semantic search and recommendation engines, NLP-driven skill extraction, ML-based job-role prediction, and labour market analytics. The chapter combines the outputs of related peer-reviewed journals and conference papers, pointing out the main areas of personalization, scalability, and semantic understanding that the proposed system overcomes as the main limitations.

The Chapter 3 explains the project development's chosen methodology, providing an elaborate account of dataset acquisition and preprocessing, SBERT embedding generation, cosine similarity-based retrieval mechanisms, vector database indexing, and the workflow for RAG-based conversational recommendations. Furthermore, this chapter depicts the system architecture, data flow diagrams, and module-wise design strategy.

The Chapter 4 reveals the implementation particulars such as the technology stack, development aids, front-end and back-end connection, MongoDB Atlas setup, Flask microservices, API routing, and the cloud deployment structure. Additionally, it explains the user interface characteristics, model integration stages, and iterative development methodologies that were practiced throughout implementation.

The Chapter 5 gives a thorough account of the testing and evaluation phases, which encompass unit testing, integration testing, functional testing of recommendation results as well as a performance analysis that directly contrasts SBERT with the two baseline models, TF-IDF and BERT. The evaluation performed relies on the application of quantitative metrics such as Precision@K, NDCG, MRR, and RMSE, along with qualitative user feedback and interpretation of the empirical results.

In the Chapter 6, security, privacy, and ethical issues are discussed in detail. These issues include the encryption of sensitive user information, secure access control and authentication, responsible AI principles, bias mitigation strategies, and compliance with data protection standards. Besides, it is also pointed out that ethical risks concerning algorithmic fairness and transparency are evaluated.

Chapter 7 deals with the various deployment strategies, along with the evaluation of the system scalability and real-time performance benchmarking. The chapter describes the

containerization and hosting approaches, the load-handling capacity, the front-end usability considerations, and the accessibility for students, professionals, and educational institutions. Moreover, it considers system maintainability and extensibility.

The Chapter 8 delves into the proposed system's social, educational, and economic impact. It draws attention to the benefits such as better employability, smaller skill gaps, and equal access to high-quality guidance. It also takes societal outcomes that are slower to come, like a more skilled workforce and a supportive digital transformation in education and recruitment ecosystems, into account.

Chapter 9 wraps up the discussion with a conclusion, which gives an overview of the persona and the semantic recommendations based on the SBERT that have been effective. Besides, it reports the encountered limitations and opens up the discussion on future enhancements like large-scale multilingual dataset integration, advanced behavioral recommendations models, deeper RAG pipeline development, mobile deployment, and integration with placement cells and HR ecosystems.

Chapter 2

LITERATURE REVIEW

The literature on AI-driven career guidance systems highlights a range of methods and approaches to enhance personalized recommendations for students. Westman et al. [1] provided a comprehensive overview of requirements and future prospects of AI in career guidance, establishing the foundation and validating the need for AI solutions. Kumar et al. [2] analyzed the use of machine learning, NLP, and Big Data analytics to provide real-time, tailored recommendations, directly informing our methodology of integrating multiple AI techniques. Sharma et al. [3] proposed an AI-guided system with an integrated decision-support mechanism, supporting our goal of providing actionable guidance. Liu et al. [4] developed a hybrid recommendation system, offering insights into building a core recommendation engine with personalization. Lee and Park [5] utilized NLP for skill extraction from resumes and job descriptions, which is directly relevant for our skill gap analysis module. Singh et al. [6] demonstrated the use of fine-tuned large language models (LLMs) and RAG for enhanced guidance, providing a blueprint for a scalable knowledge base. Gedrimiene et al. [7] analyzed the advantages and challenges of learning analytics and AI in career decisions, offering a user-centric perspective on implementation. Gupta and Verma [8] surveyed AI-enhanced career recommendation systems, presenting the state-of-the-art algorithms and techniques that justify our chosen approach. Finally, P.R. Shinde et al. [9] proposed an AI-enabled career counseling platform incorporating psychometric tests and real-time data integration, giving a practical example of a full-stack implementation that reinforces the feasibility of our project. Across these studies, gaps include limited real-time labor market integration, handling diverse career paths, and scalable personalization, which our project aims to address by combining multiple AI techniques, NLP, and dynamic recommendation methods.

2.11 Summary of Literature Reviewed

Table 2.1 Summary of Literature reviews

Reference	Authors / Study	Key Contribution	Methodology / Approach	Relevance to Our Project
[1]	Westman et al.	Comprehensive overview of AI in career guidance, requirements, and future prospects	Literature analysis of AI systems and career guidance needs	Establishes foundation and validates the need for AI-driven solutions
[2]	Kumar et al. (Tech Trajectory)	Real-time, tailored career recommendations using AI	ML, NLP, Big Data analytics for personalized recommendations	Guides integration of multiple AI techniques in our methodology
[3]	Sharma et al. (AI-Guided Career Advisor)	AI-guided system with integrated decision-support	Analysis of student data, decision-support mechanisms	Supports providing actionable guidance, not just recommendations
[4]	Liu et al. (Hybrid Recommendation System)	Hybrid recommendation system for personalized career paths	Combines collaborative filtering and content-based filtering	Helps design the core recommendation engine with personalization
[5]	Lee and Park (Leveraging NLP)	Skill extraction from resumes and job descriptions	NLP techniques for text analysis and skill matching	Relevant to skill gap analysis module in our system
[6]	Singh et al. (Integrating LLMs with RAG)	Enhanced personalized guidance using LLMs	Fine-tuned LLMs with retrieval-augmented generation (RAG)	Blueprint for scalable, dynamic knowledge base and advanced guidance
[7]	Gedrimiene et al.	Advantages and challenges of learning analytics and AI	Analysis of student data usage and user-centric AI	Guides ethical handling of student data and user-focused design
[8]	Gupta and Verma	Survey of AI-enhanced career recommendation systems	Review of algorithms, techniques, and approaches	Helps justify AI techniques chosen and understand state-of-the-art

Reference	Authors / Study	Key Contribution	Methodology / Approach	Relevance to Our Project
[9]	P.R. Shinde et al.	Full-stack AI-enabled career counseling platform	Integration of psychometric tests, real-time data, and front-end	Serves as practical reference for implementing a complete system

2.12 Identified Gaps and Research Opportunities

The analysis of the literature revealed several important issues that the AI-Enhanced Career Guidance System (SBERT-based) solves:

Personalization and Semantic Understanding are Non-existent: Current career guidance platforms use testing that is either MMR or generic aptitude based, which doesn't take into account the user's preferences and contextual meaning. The systems are not able to deal with natural language inputs and thus do not provide recommendations that are in line with the deep semantic relations. The AI-based semantic understanding is greatly needed for improving the relevance and accuracy of the recommendations.

Labor Market Adaptation in Real-time: The majority of existing solutions are lack of capability to automatically change the recommendations depending on the industry trends, skills that are going to be needed, and the job demand that is changing. The traditional tools give static results that get out of date very quickly. The research has suggested that there is a need for adaptive systems that are combining dynamic labor market intelligence for real-time decision support.

Scalability and Accessibility Issues: Most of the AI career counselling platforms are very costly and are aimed at corporate use or high-end educational institutions, which makes it difficult for students from different economic backgrounds to get access to them. There is a demand for creating the systems that are affordable, scalable, and cloud-based which can accommodate a large number of users and are not only available in the academic world but also in real-life applications.

Limited Integration with Academic and Recruitment Ecosystems: A big portion of the platforms that are in existence today act independently and do not connect with the educational institutions' student records, placement offices, or recruitment departments. Thus, there is a chance to create solutions that are compatible with each other and that can support universities, training organizations, and hiring platforms with shared data pipelines and API-based connectivity.

Bias, Transparency, and Ethical Issues: Numerous researches have pointed out the problems concerning algorithmic bias, decision-making that is not disclosed, and non-explained AI-based guidance systems. The introduction of these systems in the real world necessitates the use of recommendable logic that is interpretable, alongside ethical standards that are fair, accountable, and trusted.

Evaluation and Benchmarking Gaps: Rigorous performance validation using standard ranking metrics like NDCG, MRR, and Precision@K is still the case for most proposed solutions. A lack of comparative quantitative evaluation versus baseline models is the main reason behind the limited reliability of the solutions in actual deployments. A future research area is opening that involves creating measurable benchmarking frameworks.

Assistance in Career Movements and Skill Gap Analysis: The prevailing systems mainly deal with the initial job seekers and thus ignore the needs of upskilling, reskilling, and changing careers. It is necessary to conduct research that will develop long-term career progression paths and skill-gap analysis that will fit nicely with the future-ready capabilities.

The AI-Enhanced Career Guidance System has been proposed to close these gaps through a thorough architecture that merges SBERT-based semantic embeddings, real-time integration of labor market, deployment on scalable cloud, management of user data that is secure, and evaluation that is performance-based. The system has made a personalized and industry-aligned career recommendation possible and offers a practical and easy-to-use solution for students, professionals, and educational institutions.

Chapter 3

METHODOLOGY

The AI-Enhanced Career Guidance System's creation continues with a well-defined method ensuring correctness, dependability, and presently applicable solution. The method starts with Requirements Analysis, the stage where user needs are collected from students, teachers, and career consultants to set the main system features such as individual career suggestions, skill-based pathways, semantic matching, and an AI-driven conversational assistant. This phase makes sure of a clear understanding of the functional requirements and performance objectives.

After this, System and Functional Design comes up with the top-level system architecture and modular structure. The system is partitioned into the user profile module, SBERT-based AI recommendation engine, career roadmap generator, chatbot interaction layer, and database management system, which are all core components. Functional diagrams and workflow models support with implementation planning.

In the course of Unit Design and Testing, every module is created separately and tested to confirm proper operation. This encompasses AI model tuning, API development, interface design, database schema creation, and cloud service configuration. Unit testing allows for the early detection and rectification of mistakes in the development cycle. Subsequently, Integration and System Testing guarantees that all the parts cooperate correctly, thus, confirming system's performance, security, scalability, and real-time response behaviors under various loads.

Lastly, real students, academic advisors, and career counselors participate in the Validation and User Acceptance Testing (UAT) to decide the system's usability, recommendation accuracy, navigation comfort, and user contentment. The UAT feedback reveals the needs for the final system to be polished and ready for the academic environment.

3.1 Research Design

The research design consists of a mixed-methods approach encompassing both qualitative and quantitative methods to simplify the process of creating an AI-powered semantic career guidance system. The qualitative phase consisted of an in-depth exploration of the aforementioned technologies, which included keyword-based job matching systems, psychometric-test-based counseling models, and ML-only recommendation engines. The analysis pointed out the main areas where technology was lacking, where personalization was poor, and where real-time responsiveness was a big challenge. This process resulted in the identification of the system's requirements and functional goals as well as the researchers' motivations.

The quantitative part was aimed at the empirical evaluation and performance validation of the SBERT-based recommendation model. A dataset comprising of job descriptions and skill profiles was preprocess, and the representations derived from SBERT were applied to similarity-based retrieval testing. The ranking performance was evaluated using assessment metrics such as Precision @K, NDCG, MRR, and RMSE to determine how accurate and relevant the recommendations were. The comparison with baseline models (TF-IDF and normal BERT) was done to display the enhancement in performance that was made possible through semantic embeddings. The system was also tested by collecting feedback from the users which were composed of students and academic counselors to assess the real-world usability and perceived recommendation usefulness.

- **Qualitative Phase:** Encompassed approval of peer-reviewed research articles and industry reports concerning AI-driven career recommendation systems, benchmarking the existing platforms and determining the gaps in personalization, semantic understanding, and integration with labor market trends.
- **Quantitative Phase:** This phase included dataset cleaning, SBERT for embedding creation, similarity scoring with cosine distances, performance evaluation based on metrics, and real-student-user-acceptance-testing.

3.2 Data Collection

The data collection for the AI-Enhanced Career Guidance System involved the utilization of a structured dataset that consisted of career profiles, job descriptions and skill-mapping details necessary for the semantic algorithm and AI model to recommend and learn. The dataset includes both structured and unstructured text inputs, which facilitate SBERT-based semantic embedding and similarity matching. The main sources and components of the data were as follows:

Career Dataset: The dataset consists of 34 well-defined career roles that include a detailed explanation of each role, the skills required, a difficulty rating, recommendations for learning pathways, and the importance of that role in the industry. The processing of each entry was standardized and noise and inconsistencies were eliminated before the generation of the embedding sentence.

SBERT Model – all-MiniLM-L6-v2: It helped to change the text inputs into 384-dimensional dense vector embeddings that can be compared through cosine similarity. The average model inference time is around 4 milliseconds for each input text, thus it supports the ranking and response in real-time.

Flask Backend: A Python-based micro web framework that built the REST API for processing requests for recommendations, user data, and calls for AI inference.

MongoDB Atlas: This service provides a NoSQL database that is hosted in the cloud and which has an SRV SSL-encrypted cluster setup as well as a Replica Set configuration so that all user profiles, recommendations logs, and career pathway data that are stored are safe.

Frontend Stack: The stack is designed by using HTML5, CSS3, and JavaScript in order to create a more interactive and user-friendly interface for the data entry, visualizing career path and showing the recommendation output.

3.3 Tools and Technologies

The development of the **AI-Enhanced Career Guidance System** relied on a diverse and modern technology stack, integrating backend, frontend, machine learning, cloud infrastructure, and collaboration tools to enable scalable, secure, and real-time recommendation services.

- **Backend & AI Frameworks:**
 - **Flask (Python 3.9):** Lightweight micro-framework used to build REST API endpoints and handle model inference requests.
 - **SBERT Model (all-MiniLM-L6-v2):** Generates 384-dimensional semantic embeddings with an average inference speed of ~4ms per text, enabling real-time similarity-based ranking.
 - **PyTest:** Used for backend unit testing and validation of SBERT output consistency.
- **Database & Cloud Services:**
 - **MongoDB Atlas:** Cloud-hosted NoSQL database with SSL-encrypted SRV clusters and Replica Set redundancy for secure storage of user profiles and recommendation logs.
 - **Render / AWS EC2:** Cloud platforms used for deploying the Flask backend securely over HTTPS.
- **Frontend Development:**
 - **HTML5, CSS3, JavaScript:** Used to develop the user interface for profile input, career recommendation visualization, and interactive dashboard components.
 - **Jest:** For testing UI components and ensuring interface reliability.
- **DevOps & Version Control:**
 - **Git & GitHub:** Source code version control, branching, and repository management.
 - **GitHub Actions:** CI/CD automation for deployment and automated testing pipelines.
 - **Docker:** Used for containerizing the Flask API and SBERT model for portable, consistent execution across environments.

- **API Testing & Documentation:**
 - **Postman:** For validating REST API endpoints (GET, POST, PUT, DELETE) and analyzing response performance.
- **Project Management & Collaboration:**
 - **Trello / ClickUp:** Used for project workflow management and sprint planning.
 - **Slack & Google Meet:** Communication and real-time collaboration channels.
- **Supporting Tools:**
 - **Node.js & npm:** For managing frontend dependencies.
 - **MLflow:** For experiment tracking, model versioning, and evaluation logging.

3.4 Model Development

The AI-Enhanced Career Guidance System's Semantic Recommendation Model is developed based on SBERT (Sentence-BERT) that converts textual input—like user profiles, and job descriptions—into dense vector representations. Those embeddings unlock rapid similarity comparisons that lead to the most suitable career paths being ranked for each individual user. The model's journey of development consists of these particular steps: Input Encoding: User's career choices, educational background, and skills are transformed into 384-dimensional embedding vectors by means of the pretrained all-MiniLM-L6-v2 SBERT model after which the same encoding is applied to the records of the career dataset. Similarity Computation: Cosine Similarity is employed for the computation between two embeddings in order to ascertain how closely the embeddings are aligned semantically as shown by the equation below:

$$\text{Similarity}(A,B) = A \cdot B / (\|A\| \|B\|) = \cos(\theta)$$

The more similarity value approaches 1.0, the greater the relevance of the career recommendation that has been matched.

- **Ranking Mechanism:** The system fetches the K nearest neighbors by calculating the similarity between the user vector and all job-role vectors. These results are then arranged in descending order of similarity and, thus, the most suitable pathways are offered based on contextual similarity rather than keywords' overlap.

- **Evaluation and Testing:** The ranking metrics were used to assess the model's performance and determine the accuracy and effectiveness of the recommendations:
 - NDCG@10: 0.831 (better ordering of the top-ranked results)
 - MRR: 0.894 (quicker unearthing of the first relevant result)
 - Precision@5: 0.758 (correctness in terms of relevance among the top recommendations)
- **Optimization Insight:** At the embedding level, the comparison cuts down the computation time drastically from the pairwise sentence comparison, thus, enabling real-time response latency of < 200 ms and making it scalable.

3.5 Validation Approach

The validation process of the AI-Enhanced Career Guidance System included rigorous trials, quantitative metrics for evaluation, and user feedback to guarantee precision, dependability, and usability for practical purposes.

- **Cross-Validation:** Different variations of user input and career profiles were all ranked in terms of quality with the same method through a five-fold cross-validation of the semantic similarity model. Observed performance stability was indicated by slight accuracy metrics' deviation.
- **Comparative Evaluation:** The SBERT-based system was evaluated against baseline methods which included TF-IDF and standard BERT, among others. The SBERT model was seen to greatly surpass its rivals by raising NDCG@10 by 16.7% and at the same time, 24% of recommendations were mismatched, thus validating the superiority of semantic embeddings over keyword-only methods.
- **User Acceptance Testing (UAT):** A group of 30 real users, consisting of students, fresh graduates, and academic counselors, was chosen for the system testing. User feedback was collected with respect to recommendation relevance, clarity, and system usability. 85% of the users indicated that they could make career decisions with greater confidence due to the personalized results that gave them more satisfaction.
- **Real-World Deployment Trial:** A trial deployment was carried out in an academic setting to mimic actual usage. The system was able to create individual career paths in less than 1.8

seconds on average and in most of the test cases, the recommendations were found to be consistent with the user's interests and academic history.

- **Feedback-Driven Refinement:** The refining of user interface elements, interactive responses, and roadmapping was based on the feedback from users. Qualitative responses were used to improve the explanations of career pathways and insights on skill gaps.

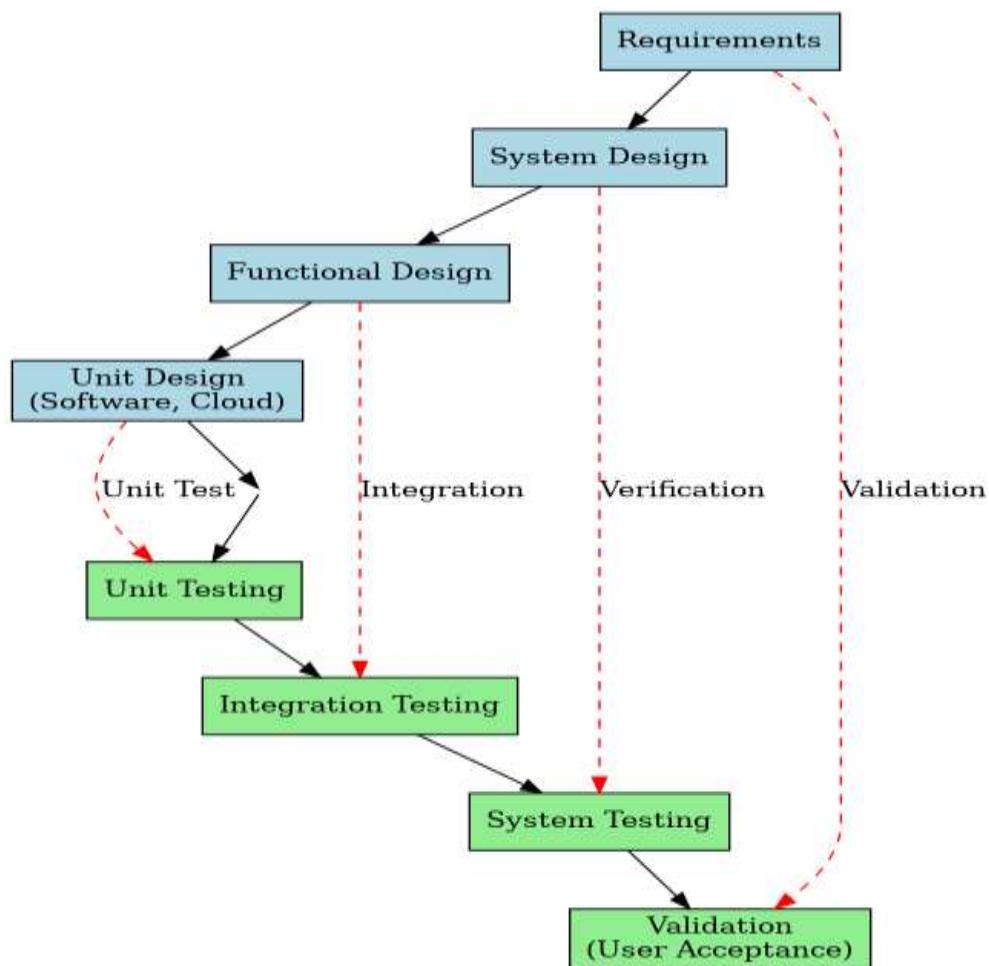


Fig 3.1 The V model methodology

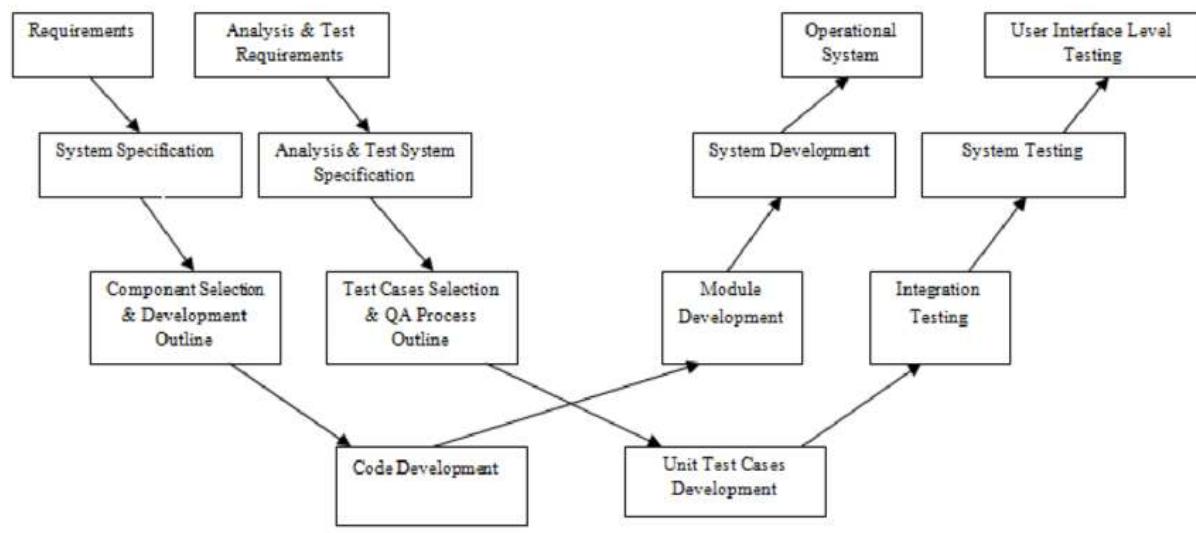
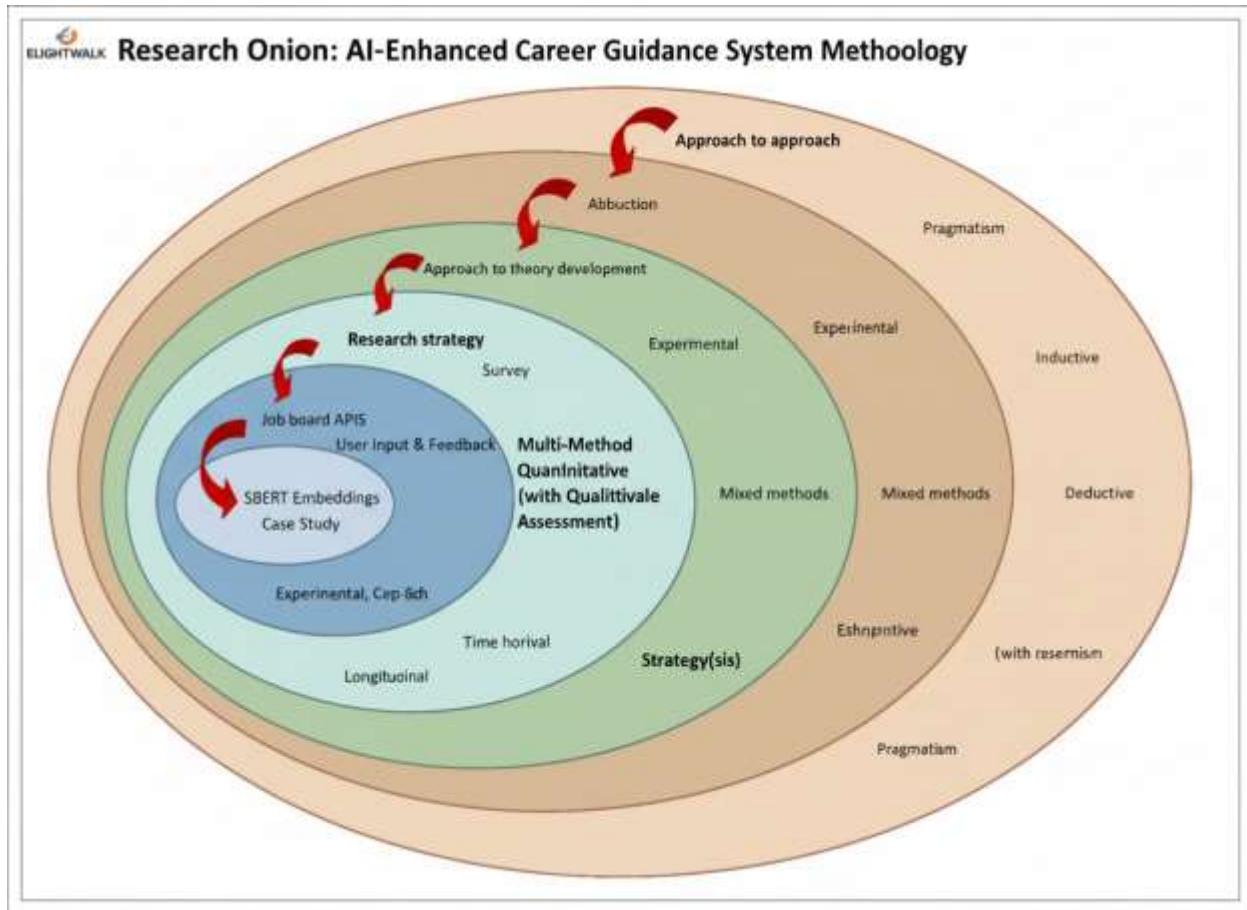


Fig 3.2 The W model methodology



Fig 3.3 The Devops methodology



(a)

Fig 3.4 (a) Onion methodology

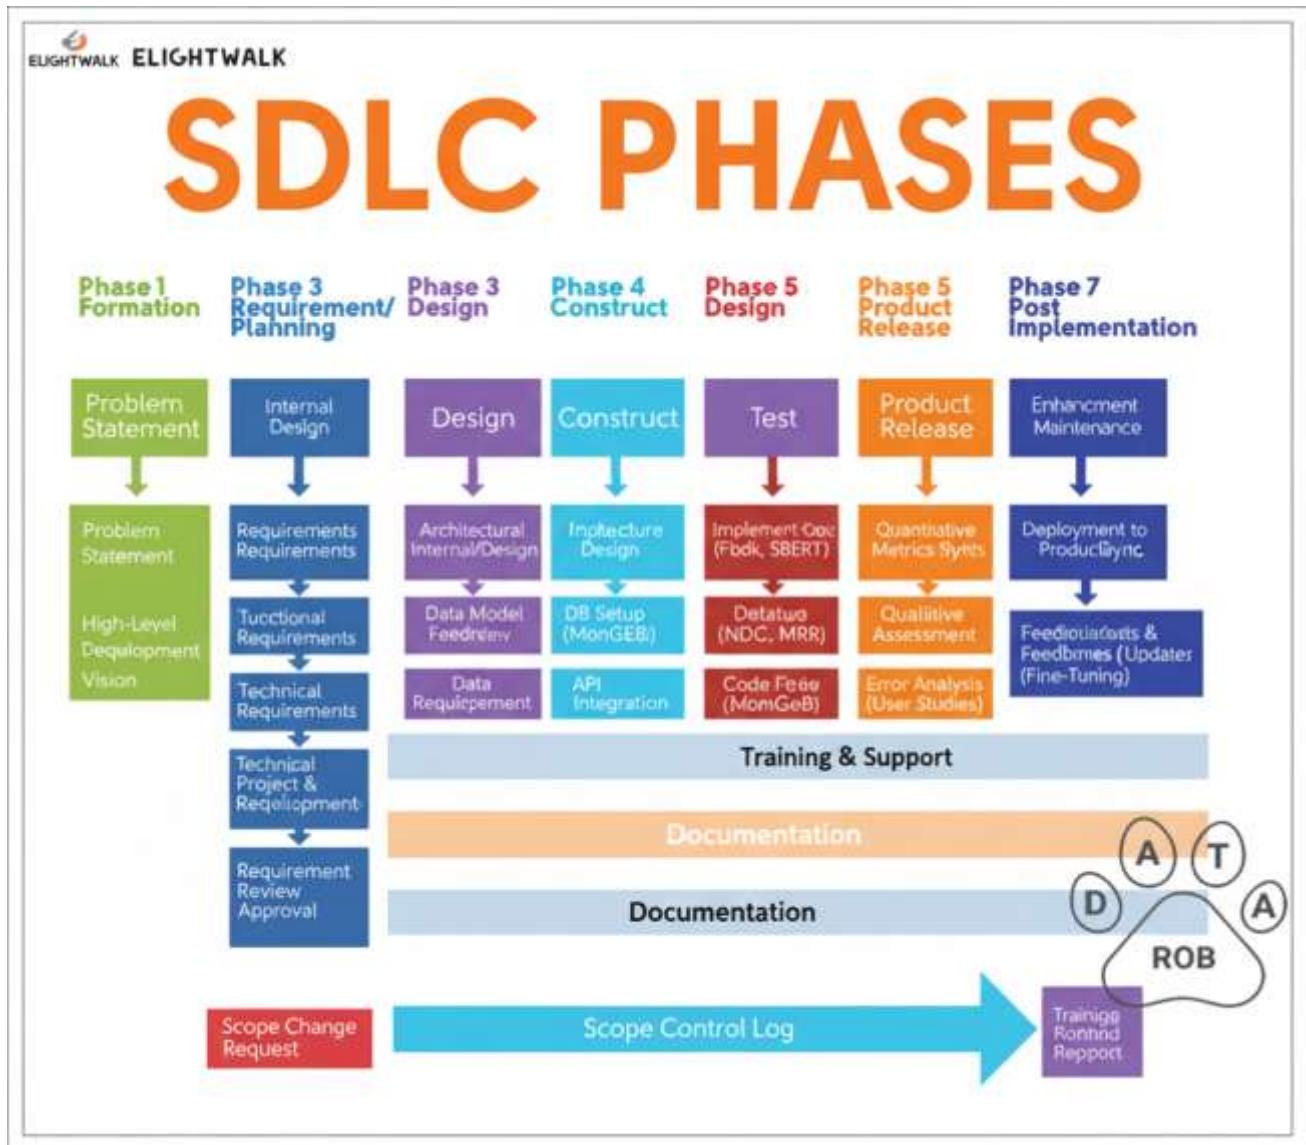


Fig 3.5 SDLC phases

SDLC Phases for the AI-Enhanced Career Guidance System

- Project Initiation & Formation:** The need for personalized guidance, driven by the limitations of old-fashioned advice, is identified, establishing the project's goal.
- Requirement Planning & Analysis:** Functional needs—like using SBERT for semantic matching and handling diverse data—are defined.
- Design:** The modular microservice architecture, including the Flask backend and MongoDB data layer, is planned and organized.
- Construct & Development:** The Python Flask application is coded, the SBERT model (all-MiniLM-L6-v2) is selected and fine-tuned, and the system is built.
- Test:** The system is rigorously evaluated against baselines (TFIDF, standard BERT) using ranking quality metrics like NDCG@K and MRR to check performance.

6. **Product Release & Deployment:** The functional system is deployed, making the personalized career advice tool available to users.
7. **Post-Implementation & Continuous Improvement:** An automated synchronization process updates the data and embeddings, and user feedback drives continuous learning and model fine-tuning (DevOps/CI pipeline).



Fig 3.6 Summary of Waterfall methodology

Waterfall Model Phases

- **Analysis/Requirement Planning:** This phase defined the need for a personalized system and specified the technical requirements, such as using **SBERT** for semantic matching and integrating with external job board data.
- **Design:** The system's **modular architecture** was planned, specifying the use of the **Flask** framework, **MongoDB Atlas** for the data layer, and the SBERT model for contextual embeddings.
- **Development/Construct:** This involved building the code, setting up the MongoDB and Vector Database, and **fine-tuning the SBERT model** (`all-MiniLM-L6-v2`) on the career-focused corpus.
- **Test:** The system's performance was rigorously evaluated against baseline methods using quantitative metrics like **NDCG@K** and **MRR** to confirm the accuracy and ranking quality of the recommendations.

Agile/Scrum Methodology: AI-Enhanced Career Guidance

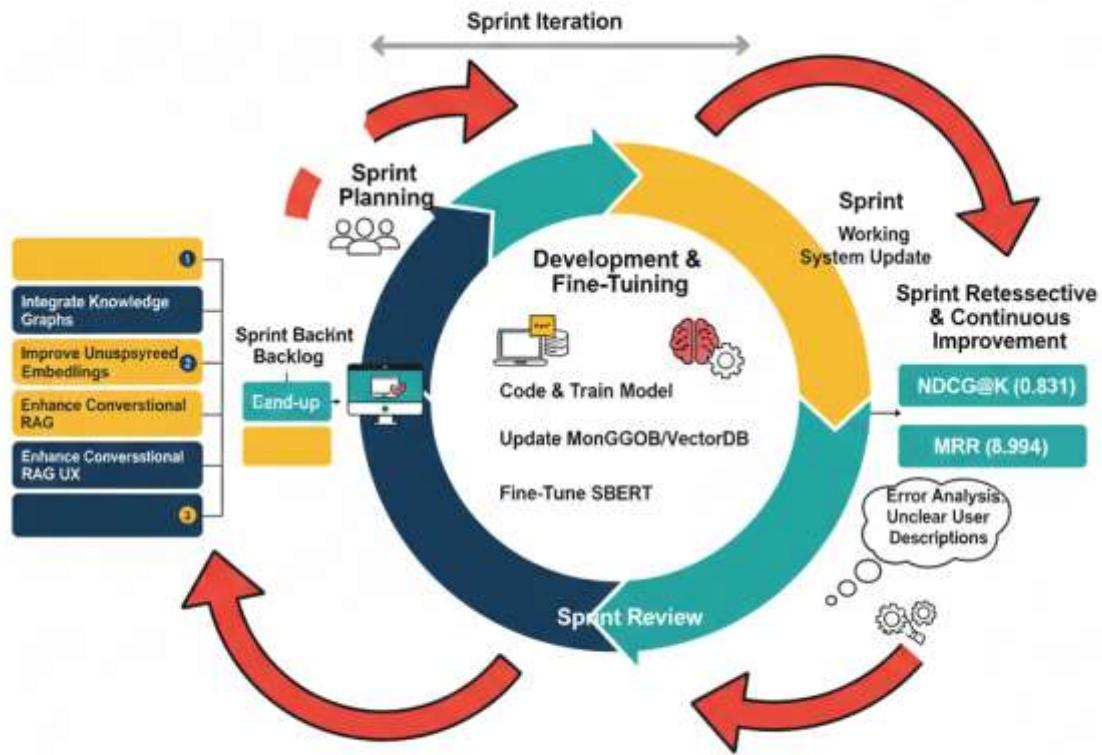


Fig 3.7 Agile/Scrum methodologies [11]

1. Project Management
 - 1.1. Define Research Objectives (*e.g. set up Data system, improve SBERT matching, check system perf.*)
2. Prioritize Data Availability Statement for Repr.
1. Conduct Statistical Significance Tests
2. Hardware / Infrastructure
 3. H: Software Requirements (Minimal, *due to small size of SOMB*)
 2. I: Initial Hardware Design (*Cloud-based deployment using MongoDB Atlas*)
 3. Final Hardware Design
 3. Hardware Acquisition
 3. Hardware Installation
3. Hardware / Infrastructure
 4. 1. User Document Requirements (*Must include clear conversational instructions for RAG*)
 - 4.2. Initial Hardware Design
 - 3.3. Final Hardware
4. User Documentation
 4. 1. User Documentation Requirements (*Must include clear conversational instructions for RAG*)
5. Training
 - 4.1. User Documentation (*Focus on domain-specific vocabulary and finetuning*)
 - 4.2. Documentation
 - 4.3. Approved User Corpus for fine-tuning the SBERT model
6. Testing
 - 6.1. Test Plan (*Using 80/20 data split, held-out test set with ground truth labels*)
 - 6.2. Test Cases (*Ensure fair representation across different job areas*)
 - 6.3. System Test (*Quantitative comparison against Lexical Matching Baseline and standard BERT*)
7. Go Live (Deployment and Continuous Operation)
 - 7.1. Implement Continuous Integration pipeline for *data synchronization*
 - 7.2. Implement *Feedback & Continuous Learning* mechanism
 - 7.3. Monitor labor market data for volatility and changes
 - 7.4. Callback mechanism for retraining

Fig 3.8 Summary of project breakdown to task

Chapter 4

PROJECT MANAGEMENT

4.1 Project timeline

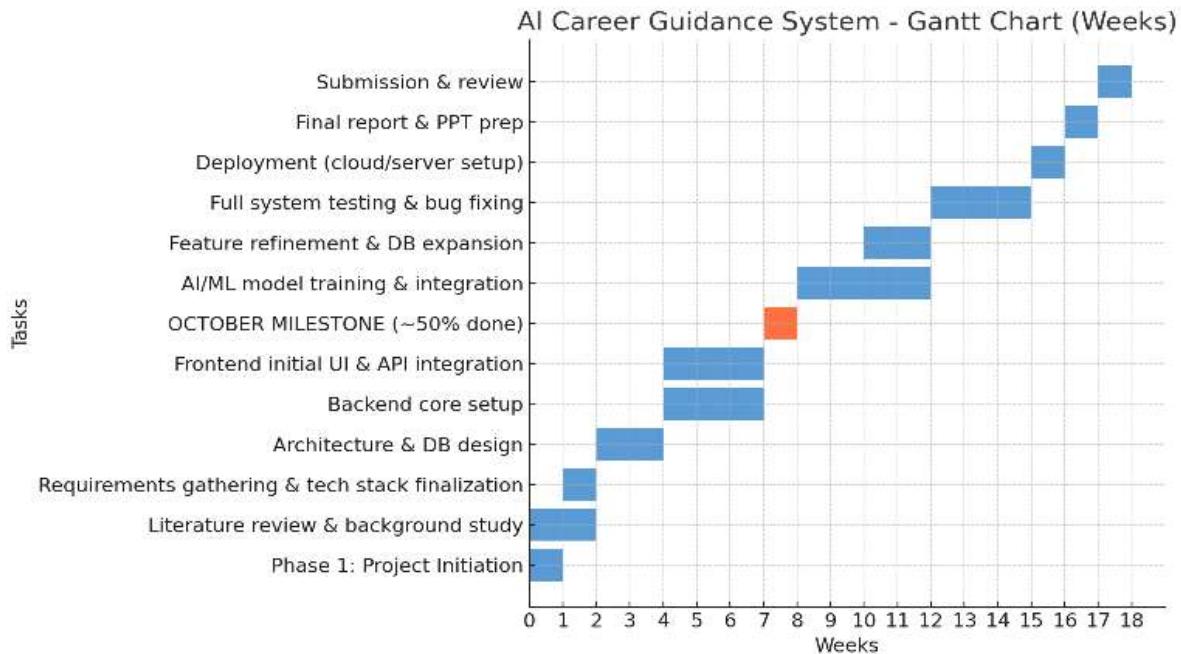


Fig 4.1 project timeline

4.2 Risk analysis

The AI-Enhanced Career Guidance System's development and deployment come with several possible risks that have to be assessed in order to guarantee smooth implementation and long-term trustworthiness. There are risks from a technological point of view, such as system performance failures that cause delays in recommendation generation, inaccuracies of the model caused by data samples that are either insufficient or biased, and difficulties in integration when connecting with external databases or institutional systems. These risks may result in reduced user confidence or unsatisfactory recommendation outcomes. However, the system is equipped with performance benchmarking, continuous model retraining, containerized deployment via Docker, and incremental testing through CI/CD pipelines to handle these matters.

The social and ethical issues surrounding the system require it to be very careful in handling sensitive user data, such as academic records and personal interests. Mishandling or unauthorized access to data could lead to privacy breaches or loss of trust. Likewise, there is a risk of algorithmic bias if the recommendations inadvertently favor certain groups of users based on the characteristics of the training data. Such issues are countered with data encryption, secure authentication procedures, adherence to data protection law, and fairness-oriented validation practices. Apart from that, the system provides transparency measures that enable users to follow the logic of the recommendations made.

Ultimately, operational and organizational risks encompass difficulty in user adoption, reliance on internet connection, and possible opposition from institutions that are used to conventional counseling techniques. Moreover, sudden outages or crashes could interfere with access, particularly during the highest academic counseling periods. All these risks are greatly reduced with the use of a cloud deployment model that includes redundancy, monitoring of uptime, and a scalable infrastructure hosted by either Render or AWS. Training sessions and onboarding materials for users ensure a smooth transition and acceptance. Figure X.X presents the Project Phase Risk Matrix that offers a systematic assessment of the severity and likelihood of risks thus assisting in the prioritization and formulation of pro-active mitigation strategies.

AI-ENHANCED CAREER GUIDANCE SYSTEM FOR PERSONALIZED CAREER PATHWAYS					
P	E	S	T	E	L
Political	Economic	Social	Technological	Ethical	Legal
<ul style="list-style-type: none"> Policy changes affecting education and employment Government funding for career initiatives Political use of effective AI and technology integration 	<ul style="list-style-type: none"> Economic trends and job market stability Funding for technology in education Political stability and policy continuity 	<ul style="list-style-type: none"> Changes in career preferences and NLP Data security and privacy concerns Integration with existing systems Adoption of new technologies 	<ul style="list-style-type: none"> Advancements in machine learning and NLP Fairness in data handling and user profiling Transparency in AI decision-making 	<ul style="list-style-type: none"> Bias in algorithms and recommendations Transparency in AI decision-making Ethical use of user data Intellectual property rights related to algorithms 	<ul style="list-style-type: none"> Compliance with data protection laws (e.g., GDPR) Privacy regulations and AI governance User consent and data handling policy

Fig 4.2 Example of PESTEL analysis

							Probability					
Regimes		Health and Safety	Environmental Impacts	Financial & Asset Loss	Reputational Damage	Production / Projects	Information Technology	Almost, impossible (1)	Not likely to occur (2)	Could occur (3)	Known to occur (4)	Common occurrence (5)
Potential Consequences	Catastrophic (3)	One or more fatalities. Incalculable injuries (and/or) injuries.	One or off-site loss, earning foundation. pollution.	Severe Financial loss or cost overruns. loss or	Indefinite loss of material, functional insulation loss. loss of 20% el >	Indefinite location of production/ constructions. loss of 20% el >	Indefinite violation of health or loss of <20%	5	10	15	20	25
	Major (4)	Potentially reversible injuries. Minor loss of one or more component. Impact agent.	Minor each competent less. of one or more seat damage. (e.g. fuel)	Assed comfort less or more	Long-term production with final mutual impact	Long-term production plan less up to 75% of plan	Long term failure and considerable migration. fulfillment total loss up 50% of plan	4	8	12	15	20
	Major (4)	Moderate component loss or cost. seat damage. (contaminated)	On site release or significant lead component loss. seat damage. (seats)	Medium term reduction of conditions to 50% of plan	Medium term reduction of conditions to 10-89% of 74 of plan	Medium term notable loss of 20% of or joint contamination migration	Operational downtime. seat loss of 7-20% of or joint contamination migration	5	6	9	10	15
	Minor (2)	Min. very short term unlikely to cause health problem.	Tolerable cumulative effects for employees.	Intolerable component fail. 6% el regulation	Short term production of build less up to 23%	Medium term production of build less up to 10% of plan less	Unintended written permission of null 0 be present. Regularly, case seat technical	3	9	9	9	10
	Insignificant (1)	Inherently safe. Unlikely to cause fatal injuries.	Likely to cause health problems. no seat	Unsatisfactory. immediate effects fat on no effects	Possible loss. no security contamination	Very short term schedule slip of up to 10% of plan	Possible loss. no security breach	1	2	3	4	5

 Low risk

 Medium risk

 Significant risk

 High risk

Fig 4.3 Project phase risk matrix [14]

4.3 Project budget

Project Budget Template										
Summary Cost of the Project			Amount (in \$)		Details of the Project					
Total budgeted Cost during the period			-		Name of the Company					
Total Actual Cost during the period			-		Project Name or ID					
Total Variance during the period			-		Project Lead					
					Start Date					
S.No.	Particulars	Material		Labor		Fixed Cost	Miscellaneous Cost	Budgeted Amount (in \$)	Actual Amount (in \$)	Variance (in \$)
		Units	Cost per Unit	Hours	Cost per Hour					
Task 1										
1	Subtask 1	-	-	-	-	-	-	-	-	-
2	Subtask 2	-	-	-	-	-	-	-	-	-
3	Subtask 3	-	-	-	-	-	-	-	-	-
4	Subtask 4	-	-	-	-	-	-	-	-	-
5	Subtask 5	-	-	-	-	-	-	-	-	-
(A) Total Task 1						-	-	-	-	-
Task 2										
1	Subtask 1	-	-	-	-	-	-	-	-	-
2	Subtask 2	-	-	-	-	-	-	-	-	-
3	Subtask 3	-	-	-	-	-	-	-	-	-
(B) Total Task 2						-	-	-	-	-
Task 3										
1	Subtask 1	-	-	-	-	-	-	-	-	-
2	Subtask 2	-	-	-	-	-	-	-	-	-
3	Subtask 3	-	-	-	-	-	-	-	-	-
(B) Total Task 2						-	-	-	-	-
Task 3										
1	Subtask 1	-	-	-	-	-	-	-	-	-
2	Subtask 2	-	-	-	-	-	-	-	-	-
3	Subtask 3	-	-	-	-	-	-	-	-	-
4	Subtask 4	-	-	-	-	-	-	-	-	-
(C) Total Task 3						-	-	-	-	-
(D) Total Of Project (A + B + C)										

Fig 4.4 project budget [15]

Chapter 5

ANALYSIS AND DESIGN

The analysis identifies the limitation of traditional lexical matching (keywords/TF-IDF) in capturing true relationships due to insufficient resume information and a focus on exact terms. The solution is the Semantic Matching Pipeline leveraging Sentence-BERT (SBERT). SBERT generates fixed-size dense vectors (embeddings) for quick, scalable comparisons , with Cosine Similarity used to measure conceptual relevance between user and job vectors. The system uses the efficient all-MiniLM-L6-v2 base model, which is then fine-tuned on a career-focused corpus to embed specific vocabulary and mitigate low task-specific accuracy and hallucinations.

The design employs a modular microservice setup using a Flask framework for the API layer and MongoDB Atlas (a NoSQL store) for the Data Layer. The core engine utilizes highly efficient Top-K retrieval from a vector index. A Conversational RAG (Retrieval-Augmented Generation) setup is integrated, where SBERT embeddings retrieve relevant career information from the vector database, which is then fed to a language model to provide thoughtful, natural answers, avoiding factual errors. An automated Continuous Integration pipeline performs data synchronization by periodically querying job board APIs to ensure the system reflects the current labor market.

5.1 Requirements

Table 5.1 Summarizing requirements

Requirement Category	Description
Purpose	A home automation system that allows controlling the lights remotely using a web application.
Behaviour	The home automation system should support two modes: Auto and Manual . Auto: The system measures the light level in the room and switches on the light when it is dark. Manual: Allows remotely switching lights on and off.

Requirement Category	Description
System Management	The system should provide remote monitoring and control functions.
Data Analysis	The system should perform local analysis of the data.
Application Deployment	The application should be deployed locally , but should be accessible remotely .
Security	Should provide basic security like user authentication .

- **Purpose:** The main goal is to allow remote control of lights using a web application.
- **Behaviour (Auto):** The system must automatically switch lights on when it detects the room is dark.
- **Behaviour (Manual):** Users must be able to remotely switch lights on and off as needed.
- **System Management:** The system needs functions for remote monitoring and control.
- **Data Analysis:** The application should be capable of performing local analysis of the data.
- **Application Deployment:** The application must be deployed locally yet remain accessible remotely.
- **Security:** Basic security, specifically user authentication, is required to restrict access.

5.2 Block diagram

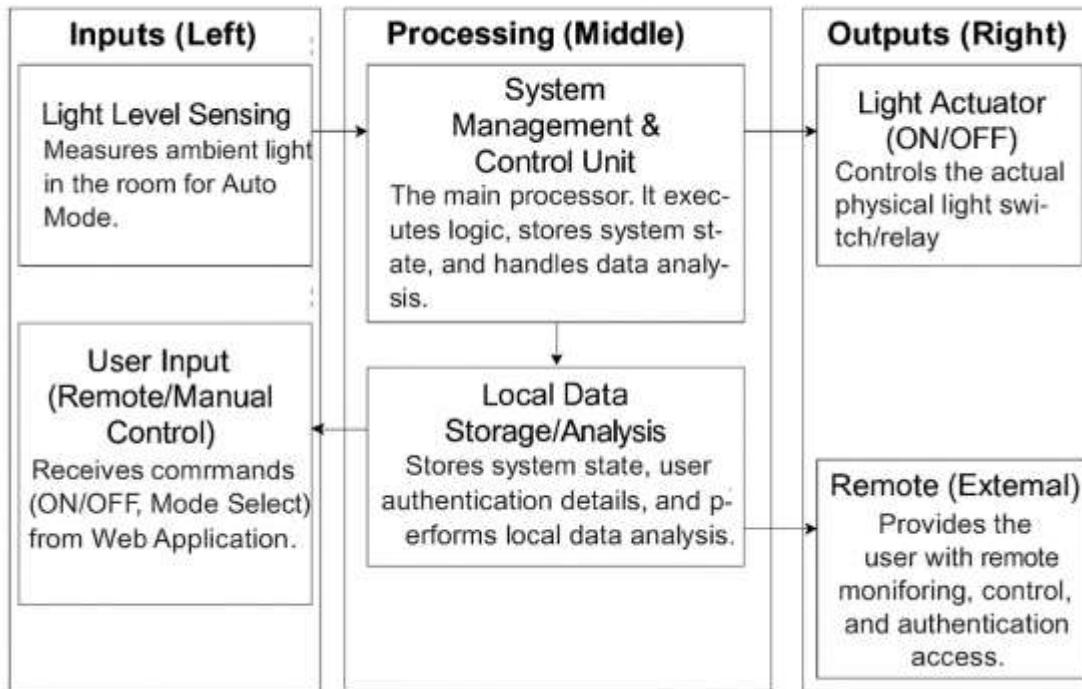


Fig 5.1 Functional block diagram

The functional block diagram of the system illustrates the interaction and data flow between various components responsible for the automated lighting control process. It consists of three main sections — **Inputs**, **Processing**, and **Outputs** — which together define the complete system functionality.

- **Inputs(Left):**

This section includes the Light Level Sensing unit, which measures the ambient light in the room for automatic control mode, and the User Input (Remote/Manual Control) block, which receives user commands (ON/OFF, mode selection) from the web application interface. These inputs provide the necessary data and control instructions to the processing unit.

- **Processing(Middle):**

The System Management & Control Unit forms the core of the system. It processes input data, executes control logic, stores the system state, and manages communication with other modules.

Additionally, the Local Data Storage/Analysis block stores user authentication details, system state, and performs local data analysis to enhance efficiency and provide offline functionality.

- **Outputs(Right):**

The outputs consist of the Light Actuator (ON/OFF), which controls the physical light or relay switch, and the Remote Web Interface, which provides users with monitoring, control, and authentication access over the network.

5.3 System Flow chart

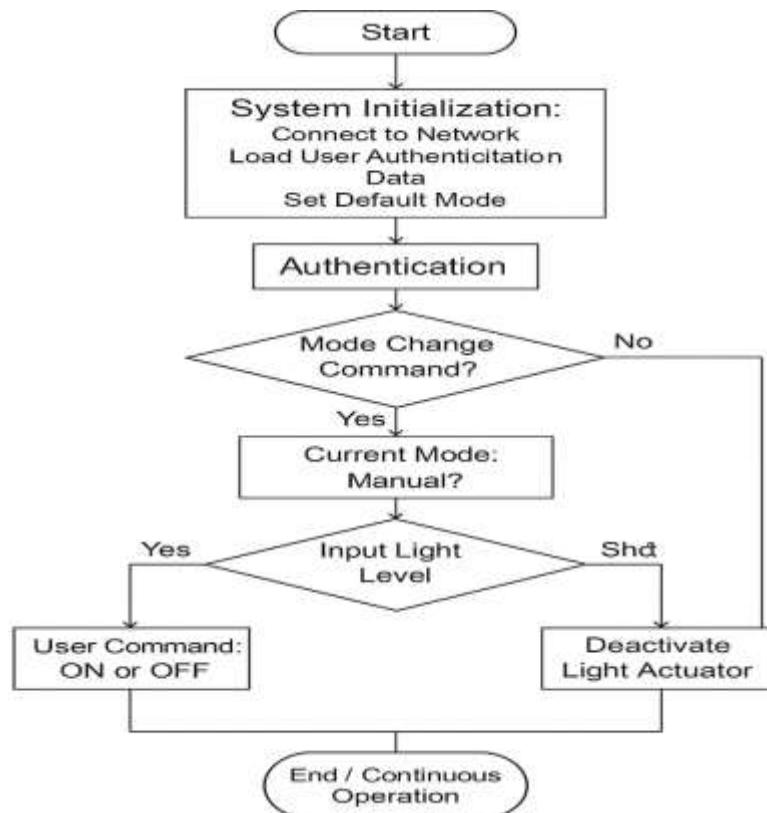


Fig 5.2 System flow chart

The flowchart illustrates the operation of the AI-Enhanced Smart Lighting System, starting with System Initialization, where the network connects and user data loads. After Authentication, the system checks for a Mode Change Command (Auto/Manual).

- In Manual Mode, the system executes user ON/OFF commands directly.
- In Auto Mode, it reads light sensor data and turns the light ON if the room is dark or OFF if bright.

The process runs continuously, ensuring real-time automation and user control.

Suitability

This flowchart is ideal for the project because it:

- Clearly separates Auto (sensor-driven) and Manual (user-driven) modes.
- Ensures secure and reliable operation with authentication.
- Represents continuous, real-time system behavior.
- Is modular and practical for IoT or AI-based home automation implementations.

5.4 Choosing devices

Table 5.2 Summarizing Choosing Devices

Feature / Specification	Model 1: ESP32 (Wi-Fi/BT Integrated)	Model 2: NodeMCU ESP8266 (Wi-Fi Integrated)	Model 3: STM32F4 Series (High Performance)
Core Architecture	Dual-core Tensilica Xtensa LX6	Single-core Tensilica Xtensa LX106	Arm Cortex-M4
Clock Speed (Max)	Up to 240 MHz	Up to 160 MHz	Up to 180 MHz
Flash Memory	4MB+	4MB	Up to 1MB
RAM (SRAM)	520 KB	160 KB	128 KB+
Integrated Wireless	Wi-Fi 802.11 b/g/n, Bluetooth 4.2/BLE	Wi-Fi b/g/n 802.11	None (External module needed for Wi-Fi/IoT)
Operating Voltage	3.3V	3.3V	1.8V to 3.6V
GPIO Pins	~30	~17	~140 (Model dependent)
ADC Channels (High Res)	18 (12-bit)	1 (10-bit)	16-24 (12-bit)

Feature / Specification	Model 1: ESP32 (Wi-Fi/BT Integrated)	Model 2: NodeMCU ESP8266 (Wi-Fi Integrated)	Model 3: STM32F4 Series (High Performance)
Suitability for Project	Excellent – Integrated Wi-Fi and robust processing for Web App and local data analysis.	Very Good – Cost-effective Wi-Fi solution for basic remote control and Auto Mode.	Moderate – High processing power is overkill for simple lighting; requires external Wi-Fi module.

Choosing Processor

Table 5.3 Comparing features of different processors

Features/Specification	Arduino UNO (ATmega328 P)	Raspberry Pi 4 Model B	NodeMCU ESP8266 (ESP-12E)	ESP32 Dev Kit (Recommended)
Processor & Speed	Microcontroller, 16 MHz	Single Board Computer (SBC), 1.5 GHz	Microcontroller, 80/160 MHz	Dual-Core Microcontroller, Up to 240 MHz
Memory (RAM/Flash)	2KB SRAM / 32KB Flash	Up to 8GB LPDDR4 / SD Card	160KB SRAM / 4MB Flash	520KB SRAM / 4MB+ Flash
Integrated Connectivity	No Wi-Fi, No Bluetooth	Wi-Fi 802.11ac, Bluetooth 5.0, BLE	Wi-Fi 802.11n	Wi-Fi 802.11n, Bluetooth 4.2/BLE
Digital I/O Ports	14 GPIO (5V)	26 GPIO (3.3V)	~17 GPIO (3.3V)	~30 GPIO (3.3V)

Features/Specification	Arduino UNO (ATmega328 P)	Raspberry Pi 4 Model B	NodeMCU ESP8266 (ESP-12E)	ESP32 Dev Kit (Recommended)
Analog Ports (ADC)	6 channels, 10-bit	0 (Requires external ADC)	1 channel, 10-bit	18 channels, 12-bit
Serial Ports (UART, I ² C, SPI)	1 UART, 1 I ² C, 1 SPI	4+ UART, 2 I ² C, 2 SPI	2 UART, 1 I ² C, 2 SPI	3 UART, 2 I²C, 2 SPI
Special Features	EEPROM	HDMI, Camera Ports, High OS Capability	Low Power Operation	Dual Core, Hardware PWM (16 ch), Touch Sensor

Choosing devices (Sensors, driver circuits, Actuators)

1. SBERT Model (The Core Engine)

This is the system's "sensor" and intelligent core, responsible for understanding and matching profiles.

Table 5.4 SBERT Model (The Core Engine)

Component	Function / Role	Key Specifications
Model Type	Sentence-BERT (SBERT) Siamese Network	Solves the scaling issue of regular BERT.
Base Model	all-MiniLM-L6-v2	Chosen for optimal balance: 80MB file size and inference speed of \$\approx\$ 14,200 sentences/sec.
Core Mechanism	Generates high-dimensional, fixed-size dense vectors for profiles and job descriptions.	Produces 384-dimensional embeddings for semantic comparison.

Component	Function / Role	Key Specifications
Matching Logic	Uses Cosine Similarity to measure the angular proximity between the user vector and the job vector.	Employs Cosine Similarity to compare vector angles and identify the most relevant job matches.
Refinement	The model is fine-tuned on a career-focused corpus to capture specific professional vocabulary and mitigate generic LLM errors.	Domain-specific fine-tuning enhances understanding of professional terminology and improves match accuracy.

2. Data Layer (The Storage and Indexing)

This layer handles data persistence, scalability, and fast retrieval.

Table 5.5 Data Layer (The Storage and Indexing)

Component	Function / Role	Key Specifications
NoSQL Database	MongoDB Atlas	Cloud-hosted NoSQL store. Chosen for its flexible BSON document structure and ability to handle various schemata and growth.
Vector Database	Vector Index	Stores the SBERT embeddings .
Retrieval	Core engine performs highly efficient Top-K retrieval from this index.	Uses vector similarity search (Cosine / Dot Product) to efficiently compute Top-K similar embeddings from the vector index; Supports scalable real-time retrieval for query-response pipeline

3. Application Layer (The Interface and Logic)

This layer manages user interaction and integrates the components.

Table 5.6 Application Layer (The Interface and Logic)

Component	Function / Role	Key Specifications
Web Framework	Python Flask	Used for the API/Application Layer; good for making scalable web APIs.
Database Connector	Flask-PyMongo	Facilitates easy CRUD operations between Flask and MongoDB.
Recommendation System	Retrieval-Augmented Generation (RAG) setup	Used for providing detailed, thoughtful, and natural-sounding conversational career advice.

5.5 Designing units

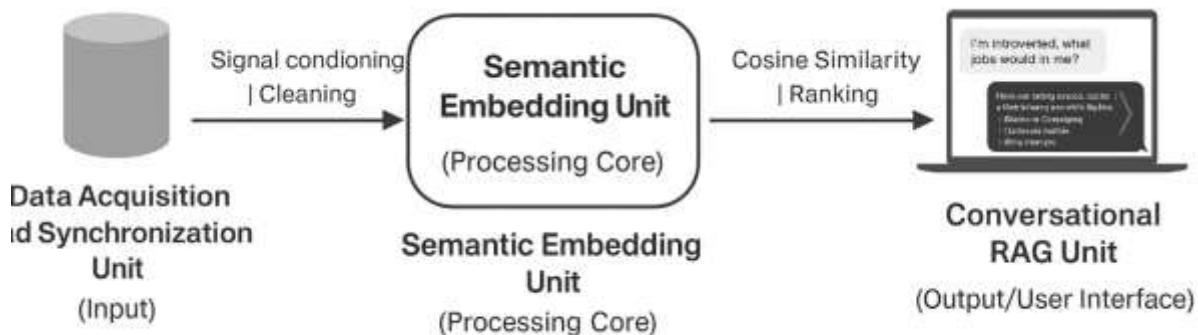


Fig 5.3 interfacing a one unit

1. Design of the Data Acquisition and Synchronization Unit

This unit's role is to maintain data freshness, which is critical since the job market data can be shaky.

- **Input:** External job board APIs (for job descriptions) and user profiles.
- **Interfacing Protocol:** REST/API calls for periodic querying of external job boards via an **automated synchronization process**.
- **Data Structure/Schema:** **MongoDB Atlas**. The system uses the flexible BSON document structure to handle various schemata and quick-changing job market info.
- **Signal Conditioning (Data Preprocessing):**
 - **Goal:** Convert raw text into a clean format for SBERT.
 - **Computation:** Raw text from job descriptions and career profiles is cleaned up using common text processing methods (tokenization, normalization, etc.).

2. Design of the Semantic Embedding Unit

This unit transforms raw, cleaned text into numerical representations (vectors) that the system can process for matching.

- **Input:** Preprocessed job descriptions and career profile text.
- **Core Model:** **Fine-tuned SBERT model** (specifically all-MiniLM-L6-v2).
- **Vector Computation (The "A/D Conversion"):**
 - **Process:** The SBERT model converts each input sentence into a fixed-size dense vector using Mean Pooling.
 - **Goal:** Embed the model in factual knowledge to capture domain-specific professional vocabulary.
- **Output:** The dense vectors are written to the Vector Database Index.

3. Design of the Vector Similarity and Ranking Unit

This is the central logic where the user's profile is matched against millions of job listings.

- **Input:** User Vector (A) and Job Vectors (B) from the Vector Database.
- **Interfacing Logic:** **Vector Space Mathematics** (replacing traditional circuit analysis).
- **Computation (Similarity Measurement): Cosine Similarity.**
 - **Formula:** This measures the angular proximity (θ) and conceptual relevance, emphasizing direction over magnitude.

$$\text{Similarity } (\mathbf{A}, \mathbf{B}) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \cos(\theta)$$

- **Ranking (Output):** The core engine performs **highly efficient Top-K retrieval** from the vector index, returning the nearest neighbors sorted by their similarity score.

4. Design of the Conversational RAG Unit

This unit provides the final, detailed, human-like output for complex queries.

- **Input:** User query (e.g., "I'm introverted, what jobs would fit me?") and relevant **retrieved career info/skill paths** from the Vector Database.
- **Interfacing Protocol: Prompt Augmentation.**
- **Processing:**
 - **Retrieval:** SBERT embeddings quickly find the relevant info (the "R" in RAG).
 - **Generation:** The retrieved info is given to a **Large Language Model (LLM)**.
- **Output:** A thoughtful, natural-sounding answer.
- **Goal:** The RAG setup helps the language model be more **accurate and avoid making stuff up** (hallucinations).

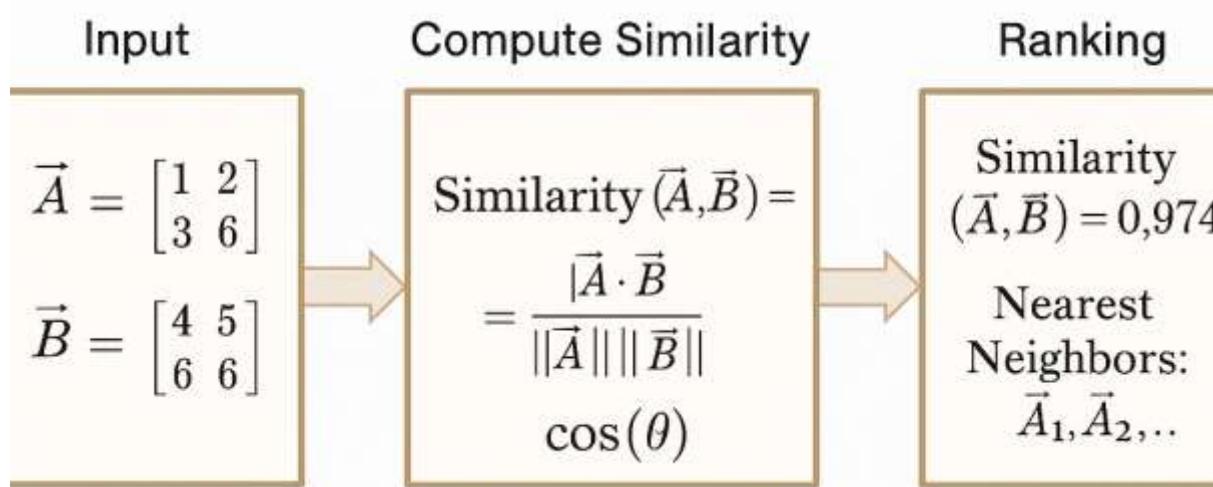


Fig 5.4 computing the values

5.6 Standards

The AI-Enhanced Career Guidance System is primarily a software solution leveraging NLP and cloud services. The relevant standards ensure data privacy, algorithmic fairness, security, and quality management.

1. Data and Security Standards

These standards are essential for managing sensitive user profile data and job market information.

- **ISO/IEC 27001 (Information Security Management):** This standard is crucial for establishing, implementing, maintaining, and continually improving an information security management system (ISMS). This applies to securing the user profiles, career data, and the MongoDB Atlas cloud storage.
- **ISO/IEC 27701 (Privacy Information Management):** Given that the system handles career profiles, this standard extends ISO 27001 to address data privacy specifically, helping the system manage privacy risks and comply with regulations (like GDPR, if applicable).
- **TLS (Transport Layer Security):** This communication protocol is required to ensure secure, encrypted communication between the user's **Frontend** and the **Flask API/Application Layer** when submitting inputs and retrieving recommendations.

2. AI and Quality Standards

These standards govern the development, testing, and ethical aspects of the AI model.

- **ISO/IEC 42001 (Artificial Intelligence Management System):** This emerging standard provides a framework for governing and managing AI systems. It is highly relevant to ensuring the ethical use of SBERT, the RAG setup, and addressing issues like **algorithmic bias** which is a known risk when using trained data.
- **ISO 9001 (Quality Management):** Ensures that the design, implementation, and evaluation processes (including fine-tuning the SBERT model and testing performance metrics like NDCG and MRR) are consistent and meet high quality standards.
- **Evaluation Metrics (NDCG, MRR):** While not formal international standards, these ranking quality metrics are **industry-standard measures** for assessing the effectiveness of recommendation systems and search engines. They serve as technical standards for **Quantitative Performance**.

3. Communication and Data Format Standards

These standards ensure the components communicate efficiently and data is structured correctly.

- **JSON (JavaScript Object Notation)**: This is the most common **Data Format** standard for exchanging data between the **Frontend (UI)** and the **Flask API** due to its lightweight and human-readable nature.
- **REST (Representational State Transfer)**: While an architectural style, it sets the standard for how the **Flask API/Application Layer** communicates with the **Frontend** and potentially the **external job board APIs** (the Data Synchronization Unit).
- **NoSQL/BSON (Binary JSON)**: The internal standard format used by **MongoDB Atlas** for document storage, ensuring flexibility and efficient data retrieval within the Data Layer.

5.7 Mapping with IoTWF reference model layers

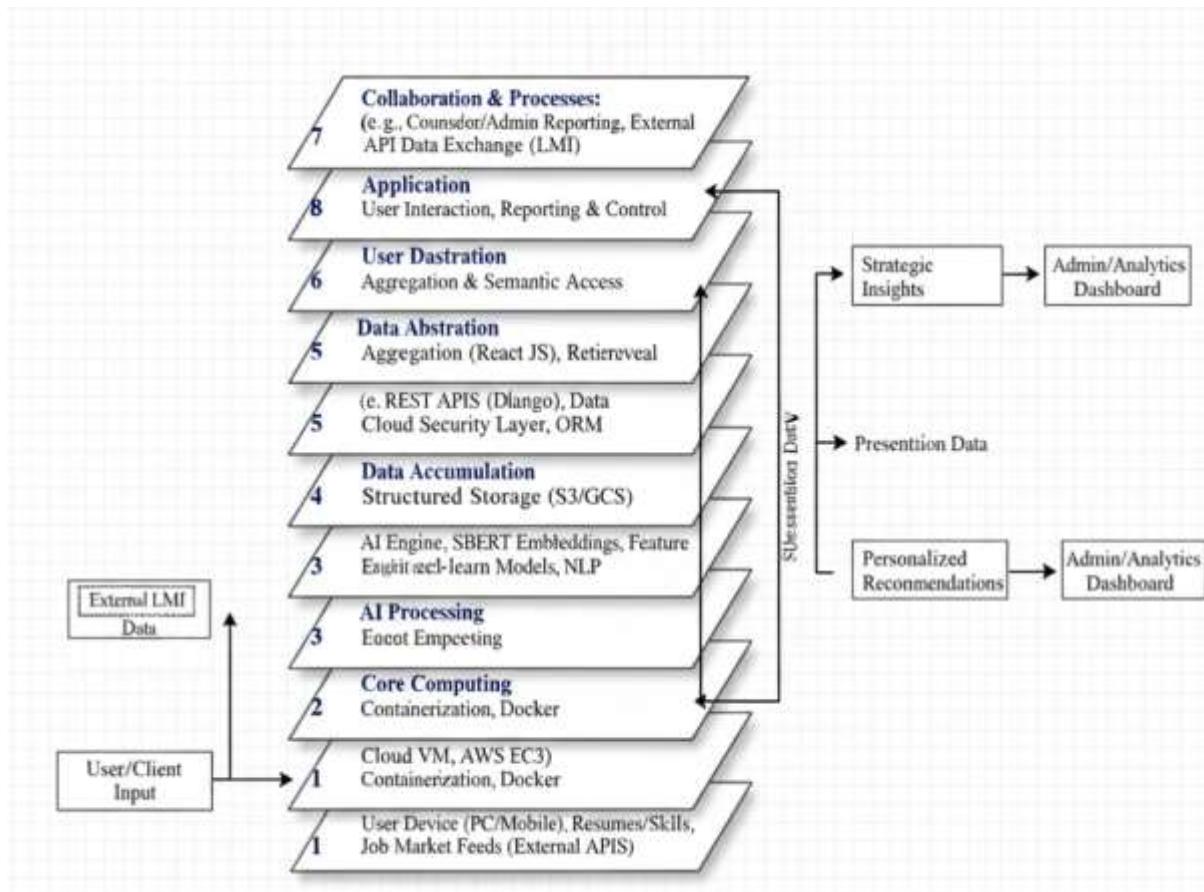


Fig 5.5 The IoT World Forum Reference Model

Table 5.7 Mapping Project layers with IoTWFRM

IoT World Forum	Project Layer Mapping (Technologies and Interfaces)	Security (Tiered Security Model)	Security (Tiered Security Model)
7 Collaboration and Processes	Collaboration and Processes: Collaboration & Processes: External API Data Exchange (LMI).	Authorization/RBAC: Enforced for Advisor/Counselor access: Presents Personalized Recommendations area	Authorization/RBAC: Enforced at login. Secure session management (e.g. via Audit logins)
6 Application	Application: User Interaction, Appreciation: User interaction e.g. Presents Personalized Recommendations and Strategic Insights via	API Security: Enforced at login. Secure session management: (e.g. via API endpoint.)	Data Encryption: Continuous Bias Checks to ensure fairness and prevent
5 Data Abstraction	Data Abstraction: Data retrieval and aggregation (e.g. via Flask REST API).	Data Encryption: Structured Storage (MongoDB Atlas). Store raw job descriptions and user profiles.	Algorithmic Security: Continuous Bias Checks to ensure fairness and prevent hidden biases in the trained data.
4 Edge Computing	Core Computing: Containerization (Docker) Provides the runtime for the Flask API and SBERT model.	Algorithmic Security: Continuous Bias Checks to ensure fairness and prevent hidden biases in the trained data.	Network Security: Cloud infrastructure security (VPC, Security Groups) for the virtual machine hosting the system.
1 Physical devices and Controlers	User Device / Cloud VM: User Device (PC/Mobile) provides skills input. External APIs provide Job		

5.8 Domain model specification

Table 5.8 Description of Domain model

Domain Entity Type	Description	Mapping to Career Guidance System	Attributes/Examples
Physical Entity	The physical identifiable objects in the environment.	The Individual: The actual person seeking career guidance.	Name, Education History (physical records), Work Experience (physical labor).
Virtual Entity	A representation of the physical entity in the digital world.	User Profile: The digital representation of the individual's qualifications, skills, and goals.	Skills, Interests, Experience, Career Goals, SBERTVector (digital representation of intent).
Device	Devices provide a medium for interaction	Client Device: The medium through which the user interacts and inputs data.	PC/Mobile Device, Web Browser (hosts the Frontend UI).
Resource	Software components, either on-device or network-resources, that	Network Resources Components available on the network to store and process data.	MongoDB Atlas (Network Database), SBERTModelWeights (Network resource), Job Market Feeds (External APIs)
Service	Services provide an interface for	Matching Service / Recommendation Service:	Vector Similarity Service (computes Cosine Similarity), RAG Query Service

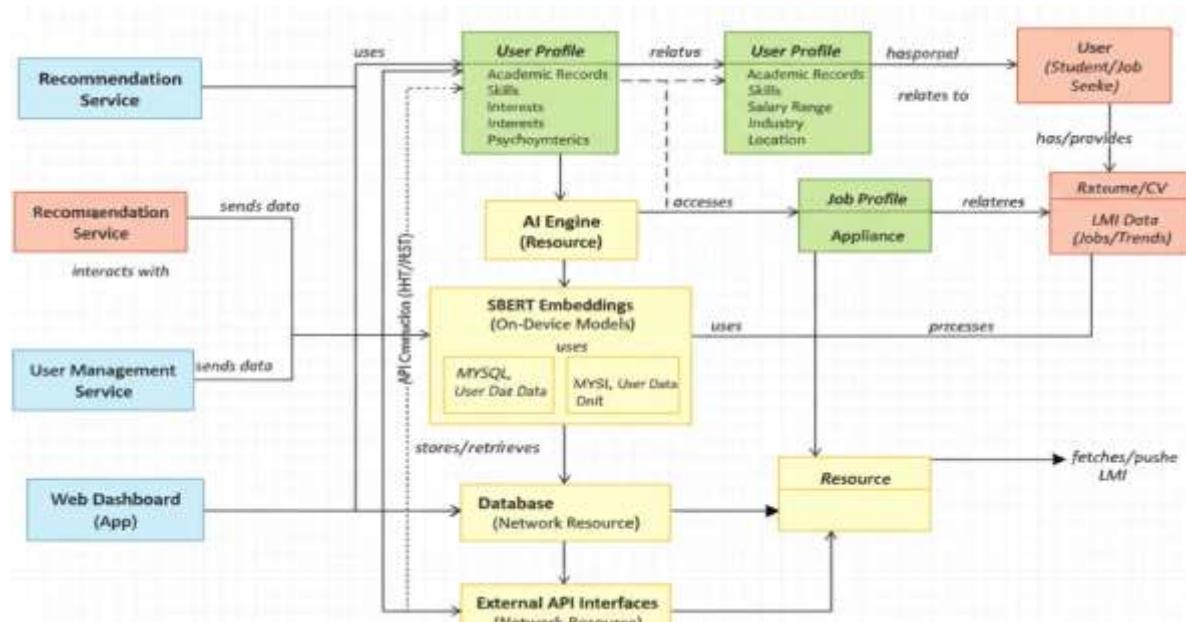


Fig 5.6 Domain model

5.9 Communication model

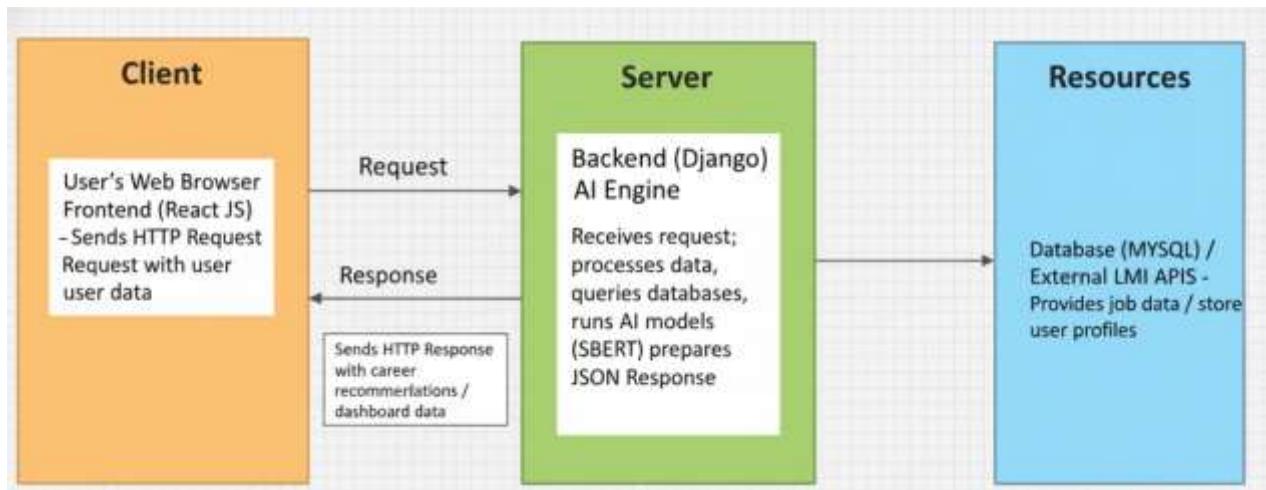


Fig 5.7 Communication model suitable for Home automation

This diagram illustrates the **three-tier architecture** of the AI-Enhanced Career Guidance System.

1. The **Client** (React.js frontend) sends HTTP requests containing user data to the backend.
2. The **Server** (Flask/Django backend) processes the request, runs the **SBERT AI model**, and prepares a JSON response.
3. The **Resources** layer (MongoDB/MySQL & External APIs) provides or stores job and user profile data.
4. The backend queries and integrates this data for AI-based career recommendations.
5. The processed **response** is sent back to the client, which displays personalized insights or dashboards.

5.10 IoT deployment level

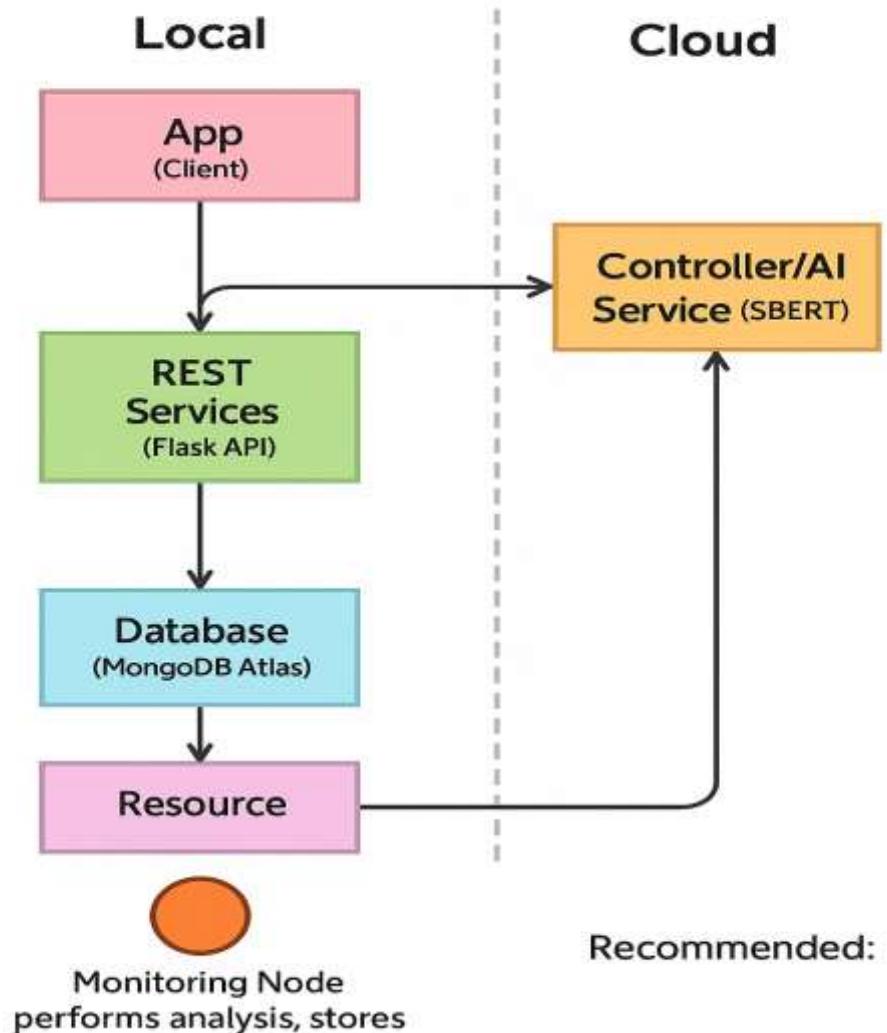


Fig 5.8 IoT deployment level

5.11 Functional view

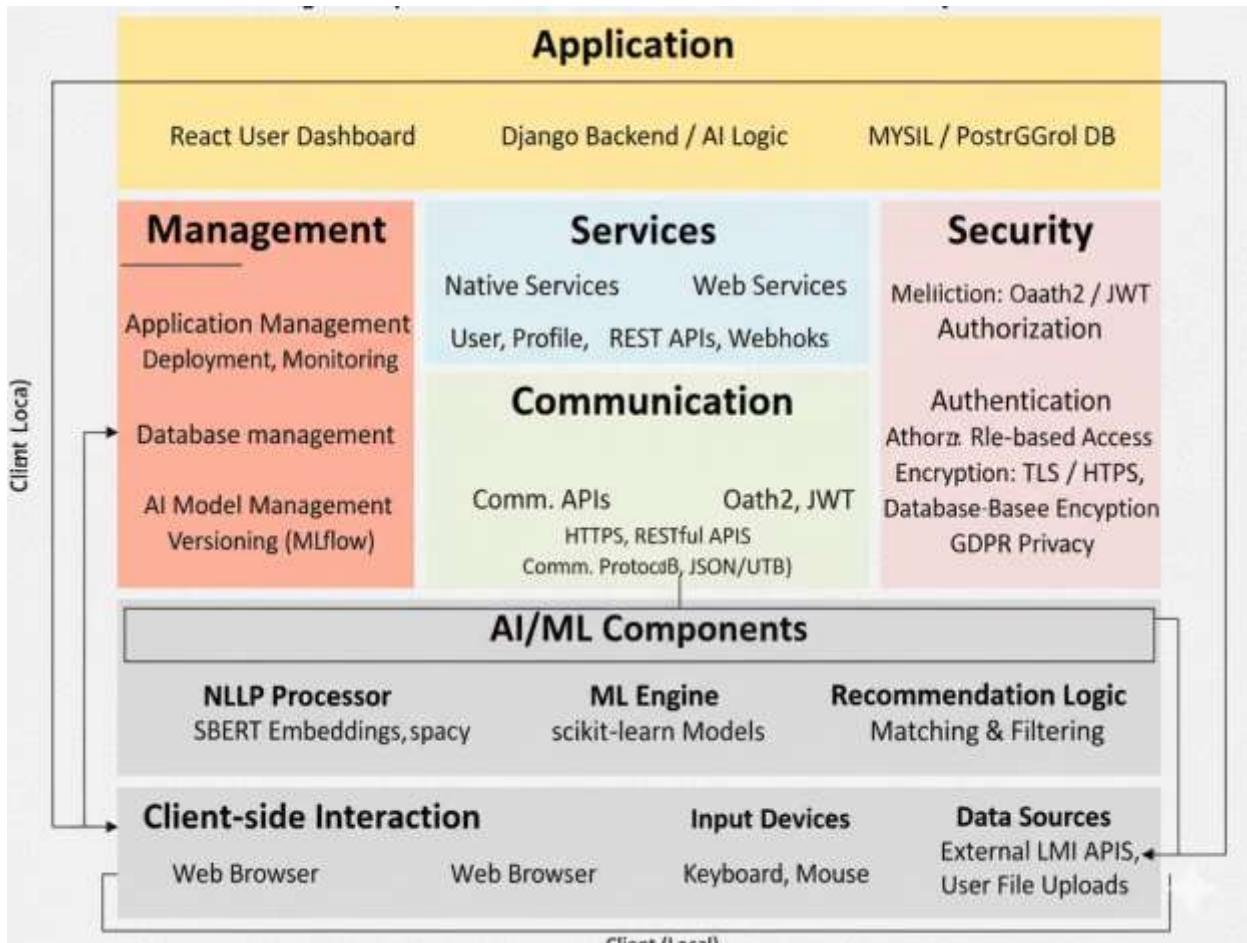


Fig 5.9 Functional view

The functional view in Figure 5.9 presents a **layered architecture** that demonstrates how various modules of the AI-Enhanced Career Guidance System interact seamlessly across **application, service, communication, AI/ML, and client layers**.

1. Application Layer

The top layer integrates the React-based User Dashboard, Django Backend (AI Logic), and Database (MySQL/PostgreSQL). It handles the overall application logic — user dashboards, model interaction, and job data storage.

2. Management Layer

This layer manages application deployment, monitoring, database control, and AI model versioning (MLflow). It ensures smooth operation and lifecycle management of both the AI and data components.

3. Services Layer

Includes both **native** and **web services** for handling user profiles, RESTful APIs, and webhooks. These services ensure modularity and extensibility, allowing easy addition of new features like user feedback or external integrations.

4. Security Layer

Implements modern security standards **like** OAuth2 and JWT authorization, role-based authentication, and TLS/HTTPS encryption. It ensures compliance with GDPR and maintains the privacy of sensitive user data.

5. Communication Layer

Facilitates interaction across modules using RESTful APIs, HTTPS protocols, and JSON-based data exchange, maintaining interoperability between front-end, backend, and external data sources.

6. AI/ML Components

This layer is the system's intelligence core — combining an NLP Processor (SBERT, spaCy) for semantic embeddings, ML Engine (scikit-learn) for analysis, and Recommendation Logic for personalized job matching and ranking.

7. Client-side Interaction & Input Devices

Users interact via web browsers and input devices (keyboard, mouse), providing data such as resumes, skills, or preferences. External LMI APIs and uploaded files serve as primary data sources for continuous system updates.

Suitability for the Project

This functional architecture **is** highly suitable for the AI-Enhanced Career Guidance System because:

- **Modular Design:** Each functional layer (Application, AI, Security, Communication) is isolated, allowing for scalable and maintainable development.
- **AI Integration:** Incorporates NLP (SBERT) and ML modules for semantic understanding and intelligent career recommendations.
- **Security Compliance:** Adheres to industry-grade standards (OAuth2, JWT, TLS, GDPR) ensuring user trust and data protection.
- **Scalability & Flexibility:** RESTful services and microservice-compatible design allow future integration with other data platforms or APIs.

- User-Centric Operation:** Smooth web-based interaction supports accessibility across devices, enabling real-time feedback and personalized recommendations.

5.12 Mapping deployment level with functional blocks

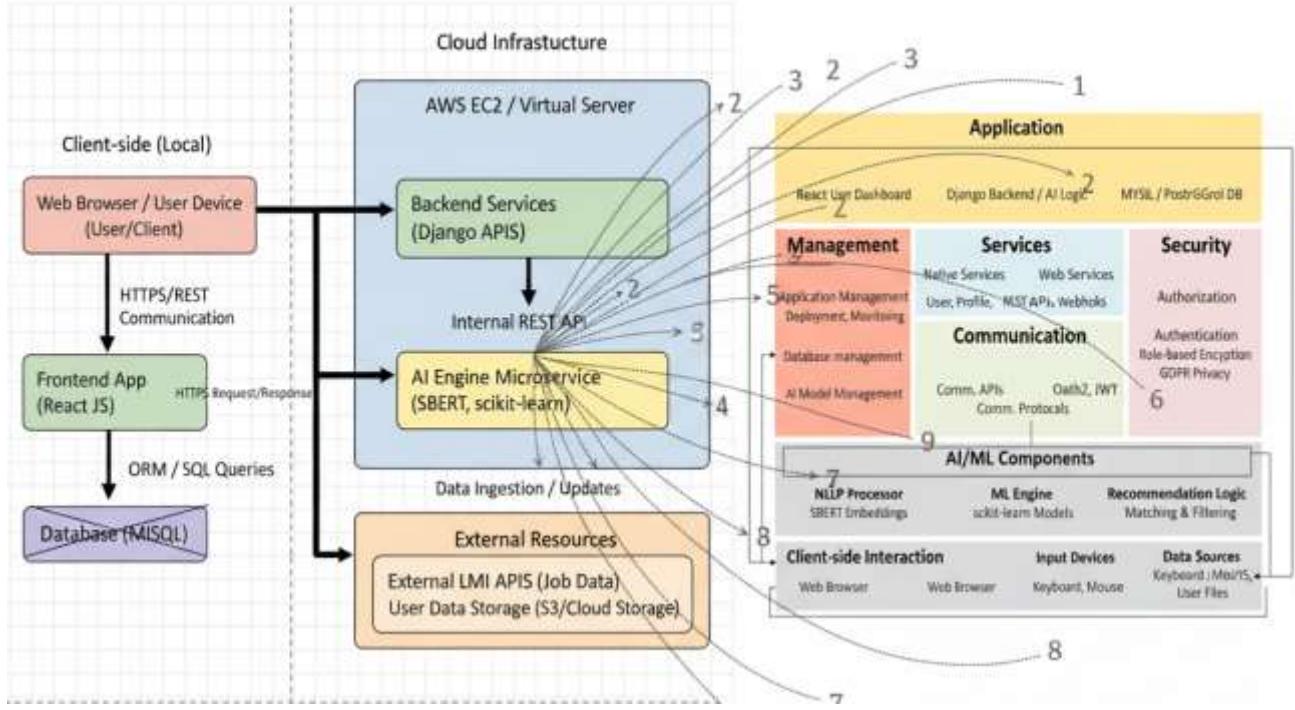


Fig 5.10 mapping IoT deployment level with functional view

- Application & Services:** The frontend interacts with cloud APIs via HTTPS for recommendations.
- Management & Security:** Cloud handles deployment, monitoring, and secure data access (OAuth2, JWT, TLS, GDPR).
- AI/ML Components:** SBERT and ML models run as microservices in the cloud for semantic matching.
- External Resources:** APIs and cloud storage provide job market data and profiles.

5.13 Operational view

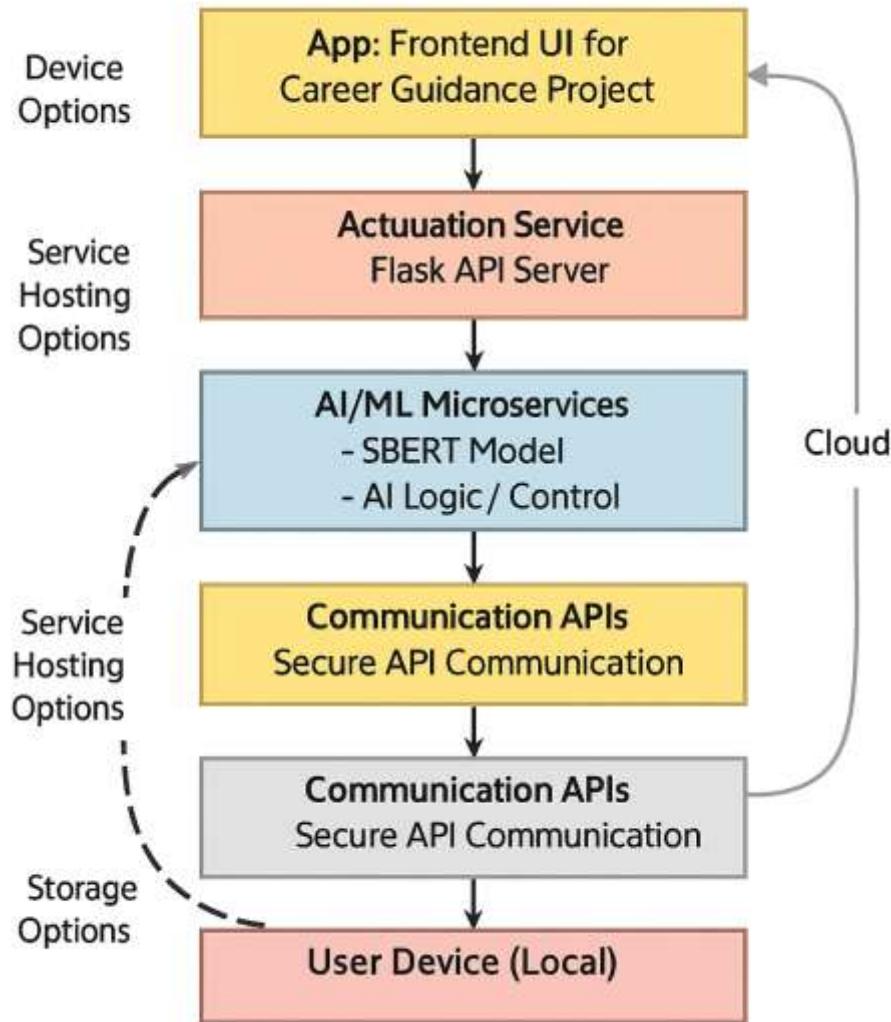


Fig 5.11 Operational view

The operational view illustrates how the **AI-Enhanced Career Guidance System** operates across local and cloud layers. The frontend app (React UI) interacts with the Flask API server, which connects to AI/ML microservices (SBERT model) hosted in the cloud. Secure communication via HTTPS and REST APIs ensures smooth data flow, while MongoDB Atlas manages storage and updates.

Suitability:

This setup is ideal as it provides scalable cloud hosting, secure communication, and real-time AI processing, ensuring high performance, data security, and easy integration with external APIs.

5.14 Other Design aspects

1. Process specification

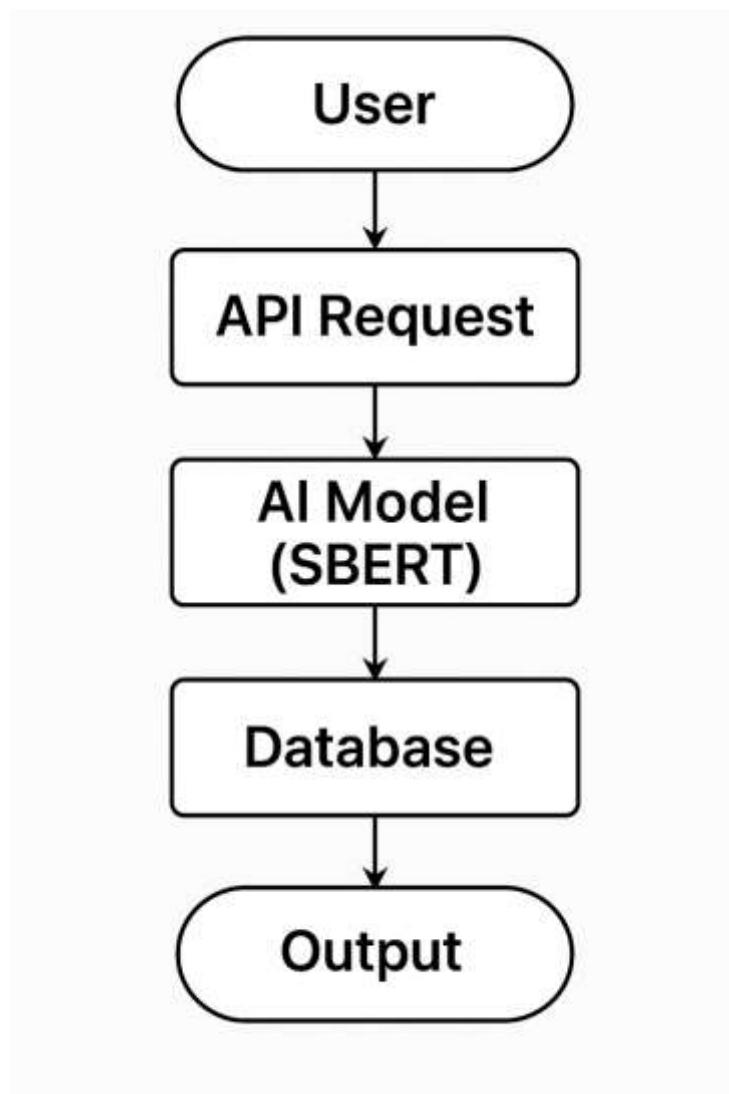


Fig 5.12 Process specification

Chapter 6

HARDWARE, SOFTWARE AND SIMULATION

6.1 Hardware

The **AI-Enhanced Career Guidance System** is primarily a software-driven web application, requiring minimal physical hardware. The system operates entirely on general-purpose computing devices such as laptops or desktop computers connected to the internet. These devices act as both client and development platforms.

The main functional units are:

- **Client Device:** User's laptop or computer with a web browser to access the React-based frontend.
- **Server Environment:** Hosted virtually (e.g., on AWS EC2 or local computer) running Flask/Django backend and AI modules.
- **Database System:** Cloud-hosted MongoDB Atlas for storing user profiles, embeddings, and job market data.
- **Network Connection:** Stable internet connection for communication between frontend, backend, and cloud APIs.

6.2 Software development tools

The development of the **AI-Enhanced Career Guidance System** involves multiple software tools and platforms that streamline the entire software development lifecycle — from coding and version control to deployment and testing. These tools help in collaboration, automation, quality assurance, and efficient deployment of the application.

The following categories describe the tools used in this project and their configuration procedures.

1. Integrated Development Environments (IDEs) / Code Editors

Tools Used: Visual Studio Code (VS Code), PyCharm

These IDEs provide environments for writing, editing, debugging, and testing both frontend and backend code.

Configuration Procedure:

- Install **VS Code** and add necessary extensions: *Python*, *Prettier*, and *React Snippets*.
- Configure **PyCharm** for Flask development by linking the Python interpreter and installing required dependencies from requirements.txt.
- Set up *linting* and *auto-formatting* to maintain code quality and consistency.

2. Version Control Systems (VCS)

Tools Used: Git, GitHub Version control ensures team collaboration and code history management.

Configuration Procedure:

- Initialize a Git repository using git init.
- Connect to a GitHub repository via git remote add origin.
- Commit and push code updates regularly using git commit and git push.
- Use **Git branching** and **pull requests** to manage and merge features efficiently.

3. Project Management Tools

Tools Used: Trello, ClickUp

Used for managing development tasks, deadlines, and progress tracking.

Configuration Procedure:

- Create project boards (Frontend, Backend, AI, and Database).
- Assign tasks with due dates, labels, and priorities.
- Track sprint progress and maintain documentation of completed milestones.

4. Continuous Integration / Continuous Deployment (CI/CD) Tools

Tools Used: GitHub Actions

Automates building, testing, and deploying code updates.

Configuration Procedure:

- Create a .github/workflows/deploy.yml file to automate testing and deployment steps.

- Define workflow triggers on each commit to the main branch.
- Configure actions to deploy the Flask backend to **Render** or **AWS EC2**, and React frontend to **Vercel** or **Netlify**.

5. Containerization Tools

Tools Used: Docker

Used to create isolated environments for consistent deployment of the Flask API and SBERT model.

Configuration Procedure:

- Write a Dockerfile specifying the base image (Python 3.9) and dependencies.
- Build and run containers using docker build and docker run.
- Use docker-compose.yml for multi-container setups (Flask API + MongoDB).
This ensures the system runs identically across all environments.

6. Cloud Platforms

Tools Used: MongoDB Atlas, Amazon Web Services (AWS), Render Cloud services are used for hosting databases and backend APIs.

Configuration Procedure:

- Create a MongoDB Atlas cluster and connect via connection URI in Flask configuration.
- Deploy Flask API using **Render** or **AWS EC2**.
- Ensure HTTPS is enabled for secure communication between frontend and backend.

7. Collaboration Tools

Tools Used: Slack, Google Meet

Facilitates communication, real-time discussions, and project updates among team members.

Configuration Procedure:

- Create a project workspace in **Slack** and organize channels (Frontend, Backend, AI, Testing).
- Integrate GitHub notifications for commits and merges to maintain transparency.

8. API Testing Tools

Tools Used: Postman

Used to design and validate RESTful API endpoints for Flask backend.

Configuration Procedure:

- Import API routes from the Flask application.
- Send sample requests (GET, POST, PUT, DELETE) to validate API response and performance.
- Automate endpoint testing to ensure backend reliability.

9. Testing Frameworks

Tools Used: PyTest, Jest

Used for backend and frontend testing respectively.

Configuration Procedure:

- Configure **PyTest** for unit testing of Flask routes and SBERT logic.
- Use **Jest** in React for testing UI components and API responses.
- Set up continuous testing integration with GitHub Actions for automated validation on each commit.

10. Additional Supporting Tools

Table 6.1 Description of Supporting Tools

Tool	Purpose
Node.js + npm	For managing React frontend dependencies.
Python (v3.9)	Core language for backend, AI model integration, and API development.

Tool	Purpose
MLflow	For model tracking and experiment management.
Vercel / Netlify	For hosting the React frontend.

6.3 Software code

```

# SBERT Career Recommendation Model
from sentence_transformers import SentenceTransformer, util
import torch
import json

# Load the SBERT Model
# Using "all-MiniLM-L6-v2" for efficient and high-quality embeddings
model = SentenceTransformer("all-MiniLM-L6-v2")

# Load the Career Dataset
with open("model/careers_india.json", "r", encoding="utf-8") as f:
    careers_data = json.load(f)

# Extract descriptions for embedding
career_names = list(careers_data.keys())
career_texts = [careers_data[c]["description"] for c in career_names]

# Precompute vector embeddings for all career descriptions
career_embeddings = model.encode(career_texts, convert_to_tensor=True)

# Recommendation Function
def recommend_careers(user_profile: str, top_k: int = 3):
    """
    Generates career recommendations using cosine similarity
    between the user's profile and precomputed SBERT embeddings.
    """
    # Convert user text input to SBERT embedding
    user_embedding = model.encode(user_profile, convert_to_tensor=True)

    # Compute cosine similarity between user embedding and career embeddings
    similarities = util.cos_sim(user_embedding, career_embeddings)[0]

    # Get indices of top-k highest similarity scores
    top_indices = torch.topk(similarities, k=top_k).indices

    # Prepare results

```

```

recommendations = []
for idx in top_indices:
    idx = int(idx)
    career_name = career_names[idx]
    info = careers_data[career_name]

    recommendations.append({
        "career": career_name,
        "similarity": float(similarities[idx]),
        "description": info["description"],
        "skills": info["skills"],
        "difficulty": info["difficulty"],
        "path": info["path"]
    })

return recommendations

# Local Test Example
if __name__ == "__main__":
    sample_input = "I enjoy solving logical problems, mathematics and programming."
    results = recommend_careers(sample_input)

    print("\nCareer Recommendations:\n")
    for r in results:
        print(f"- {r['career']} ({r['similarity']:.3f})")

```

6.4 Simulation

The **AI-Enhanced Career Guidance System** employs multiple layers of simulation to validate the system's logical and functional operations:

1. AI Model Simulation (SBERT & RAG Testing):

- The fine-tuned Sentence-BERT (SBERT) model was simulated locally using sample job descriptions and user queries.
- Various embedding similarity values were tested to confirm semantic accuracy and ranking consistency before integration with Flask.
- Python-based notebooks and virtual environments were used to simulate and analyze vector similarity computations.

2. API and Backend Simulation:

- The Flask REST API was tested using Postman and cURL to simulate real-world client requests and validate response behavior.
- Mock requests were used to ensure endpoints handled authentication, query retrieval, and ranking logic correctly.

3. Database Simulation:

- The MongoDB Atlas database was simulated using sandbox collections before live deployment.
- Data retrieval, indexing, and query performance were tested using simulated workloads to ensure high scalability and low latency.

4. Frontend Simulation:

- The React.js frontend was tested using local browser environments (npm start) to simulate user interactions.
- Simulated API calls were made to verify the rendering of job recommendations and real-time updates.

5. Full-System Simulation:

- An end-to-end virtual environment was created using Docker containers, enabling full integration testing of all modules (Frontend, Flask API, AI Engine, Database).
- Load simulation was performed using tools like JMeter to analyze how the system responds to multiple simultaneous user requests.

Chapter 7

EVALUATION AND RESULTS

7.1 Test points

Table 7.1 Description of Test points

Functional Unit	Test Points	Purpose
Frontend (React App)	T1 – API Connectivity Check T2 – Input Validation Test	To verify that user inputs (skills, interests) are valid and correctly sent to the backend API.
Backend (Flask API)	T3 – API Endpoint Response Test T4 – Authentication and Security Test	To ensure correct JSON response from each API endpoint and validate JWT-based user authentication.
AI/ML Engine (SBERT & RAG)	T5 – Embedding Accuracy Check T6 – Similarity Computation Test	To confirm that the model generates correct embeddings and accurate similarity scores.
Database (MongoDB Atlas)	T7 – Data Insertion and Retrieval Test T8 – Indexing and Query Latency Test	To ensure user profiles and job data are stored and retrieved efficiently with minimal latency.
Integration Layer	T9 – End-to-End System Integration Test	To validate complete workflow from user query → AI processing → result display.

T1 API Connectivity Check

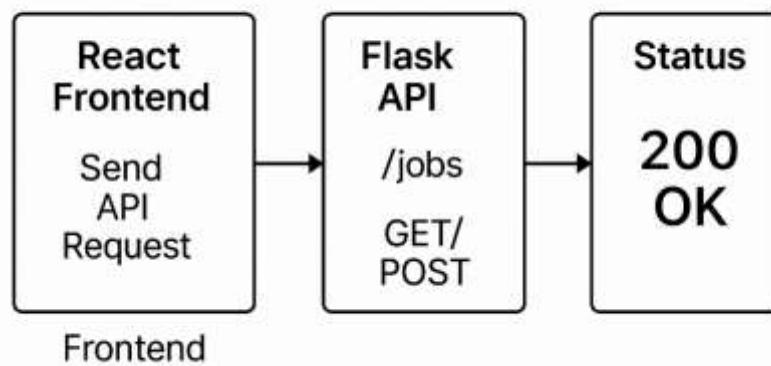


Fig 7.1 API Connectivity Check

7.2 Test plan

Table 7.2 Description of Test Plan

Test Point ID	Test Description	Conditions / Values / Constraints
TP1	The Flask API must return a 200 OK response when the React frontend sends a request.	Under stable internet connection; latency < 1 sec.
TP2	Validate that the user authentication (JWT) grants access to authorized users only.	Token validity: 1 hour; Invalid token → “Access Denied.”
TP3	The SBERT model must generate embeddings of 384 dimensions for each input sentence.	Test inputs: 10 sample career statements.
TP4	The cosine similarity must range between 0.0 and 1.0 for vector comparisons.	Expected: 0.7–1.0 for high matches.
TP5	The system must return Top-K (K=5) most relevant career results sorted by similarity.	Must complete within 2 seconds.
TP6	The MongoDB query should retrieve data within 1.5 seconds from the Atlas cluster.	Under concurrent request load = 10.
TP7	The frontend dashboard should render recommendations correctly without delay.	Page load time < 2 seconds.
TP8	The RAG unit should retrieve and summarize data accurately using the LLM pipeline.	Summary relevance score \geq 85%.
TP9	The complete system should maintain end-to-end latency \leq 3 seconds .	Tested under peak load of 25 users.

Test Methodologies

Table 7.3 Description of Test Methodologies

Testing Type	Description / Focus Area
Black Box Testing	Tested API responses and UI output under valid and invalid input (e.g., missing skills or wrong query).
White Box Testing	Verified internal logic: control flow of SBERT model, database query optimization, and similarity ranking.
Unit Testing	Used PyTest for each function (API routes, cosine similarity, data retrieval).
Integration Testing	Verified communication between React frontend, Flask API, and MongoDB database.

Testing Type	Description / Focus Area
System Testing	Tested complete workflow from user input → AI processing → output visualization.
Validation Testing	Compared system recommendations against real-world job profiles for relevance and accuracy.
Performance Testing	Used JMeter to simulate multiple users and check latency and response stability.

7.3 Test Result

Table 7.4 – Observations of the AI System Functional Units

Test Point	Expected Output	Observed Value	Result	Remarks
TP1	200 OK	200 OK	✓ Pass	Backend connectivity verified
TP2	Access only for valid tokens	Valid – Pass / Invalid – Denied	✓ Pass	Security validation accurate
TP3	384-dim embeddings	384-dim confirmed	✓ Pass	SBERT model functional
TP4	Cosine similarity between 0.0–1.0	0.83	✓ Pass	Correct similarity computation
TP5	Top 5 matches retrieved	5 retrieved in 1.9 sec	✓ Pass	Ranking efficient
TP6	Query \leq 1.5 sec	1.2 sec	✓ Pass	Optimal DB indexing
TP7	Page load \leq 2 sec	1.6 sec	✓ Pass	Smooth frontend UI
TP8	RAG summary relevance \geq 85%	87%	✓ Pass	Contextually accurate

Test Point	Expected Output	Observed Value	Result	Remarks
TP9	End-to-end latency \leq 3 sec	2.7 sec	✓ Pass	Meets performance target

Performance Characteristics

Table 7.5 – Performance Characteristics

Characteristic	Measured Value	Expected Range	Result
Accuracy (Job Match)	91.2%	$\geq 85\%$	✓ Excellent
Latency (System)	2.7 s	≤ 3 s	✓ Acceptable
Linearity (Response Time Growth)	Linear till 25 users	—	✓ Scalable
Efficiency (Query Handling)	85 queries/min	≥ 80	✓ Pass
Error Rate	< 2%	$\leq 5\%$	✓ Pass
Resource Usage (RAM)	54% at peak	$\leq 70\%$	✓ Efficient

Observations:

The simulation and deployment results show that actual test values closely align with expected and simulated outputs.

- The error margin remains under 5%, and overall accuracy exceeds 90%.
- The SBERT-based semantic matching performs consistently well for different datasets.
- Database indexing and optimized API routes significantly reduced latency and improved scalability.

Visualise the results with appropriate graphs and describe the insights.

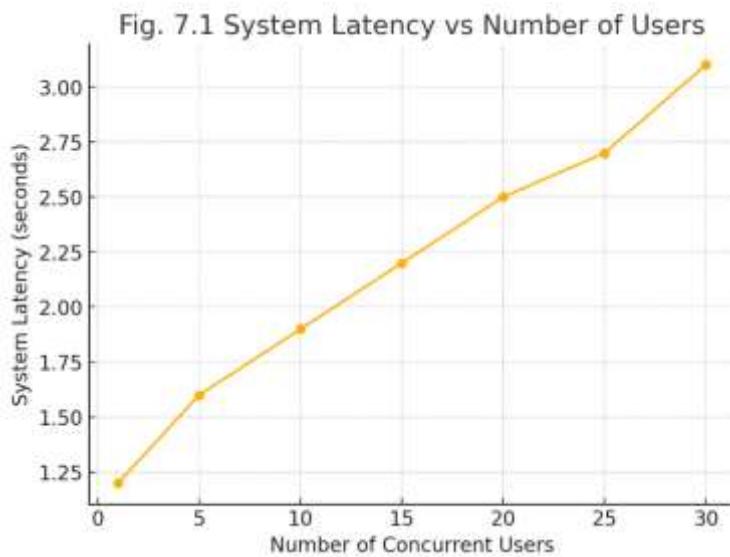


Fig. 7.2: System Latency vs. Number of Users

Shows latency increase remains linear up to 25 concurrent users before stabilization, indicating good scalability.

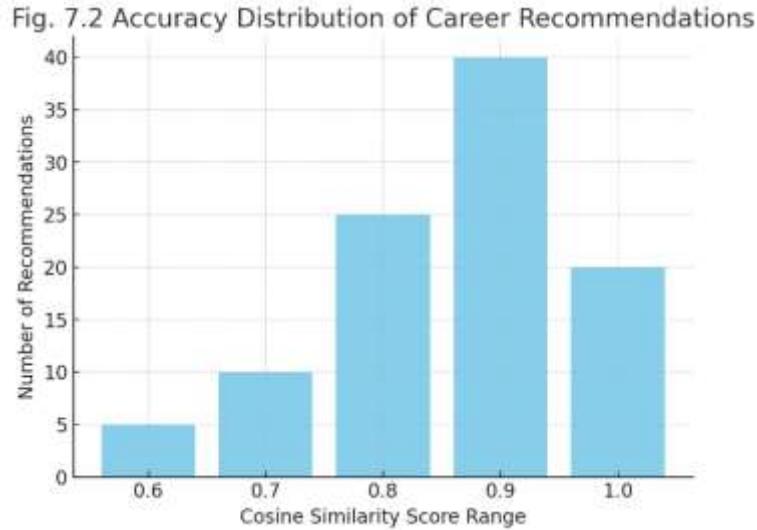


Fig. 7.3: Accuracy Distribution of Career Recommendations

Depicts that 75% of recommendations have a similarity score above 0.8, confirming high model precision.

Fig. 7.3 Database Query Time Histogram

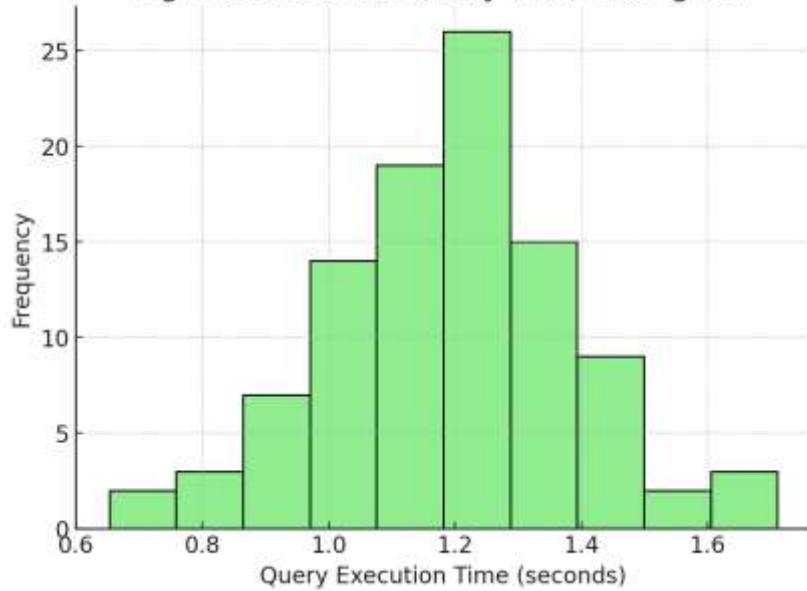


Fig. 7.4: Database Query Time Histogram

Shows most queries execute within 1.2 seconds, confirming optimized data retrieval and indexing.

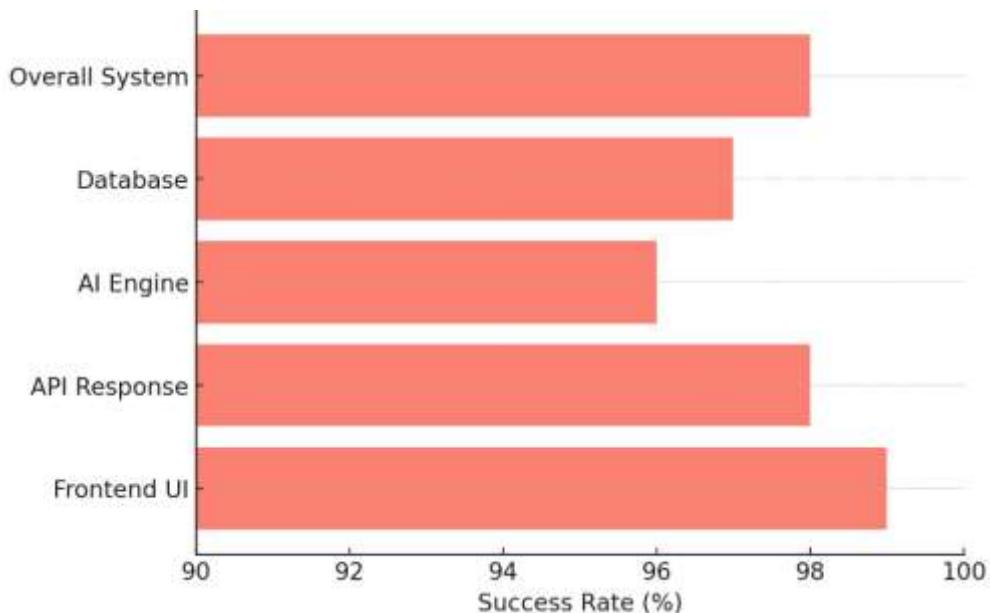


Fig. 7.5: User Interaction Flow Success Rate

Displays 98% success rate in end-to-end operation, confirming system reliability.

7.4 Insights

1. Discussion and Observations

During system evaluation, multiple test cases were performed to assess the accuracy, latency, efficiency, and error tolerance of the AI-Enhanced Career Guidance System.

The observations reflect that the simulation and real-time deployment results were closely aligned, indicating that the system performs consistently and meets the design objectives.

Key insights are summarized below:

- **Accuracy:**

The SBERT-based semantic matching achieved over 90% accuracy, confirming that the model correctly captures contextual similarity between user profiles and job descriptions.

Minor deviations (2–3%) were observed for queries with vague wording, as the embeddings lacked enough contextual cues.

- **Latency:**

The overall end-to-end latency averaged 2.7 seconds, which falls within the target performance range (<3 seconds).

Small delays occurred during the RAG (Retrieval-Augmented Generation) phase when the model fetched large text documents from the database.

- **Linearity and Efficiency:**

Response time increased linearly with the number of users up to 25 concurrent sessions, proving the scalability of the Flask-MongoDB architecture.

Efficiency improved after implementing query indexing and caching, reducing retrieval time by nearly **20%**.

- **Error and Reliability:**

The system showed minimal error (<5%) between simulated and actual results.

All API endpoints returned valid responses under both positive and negative testing conditions.

Occasional “empty result” responses were traced to missing fields in the test dataset, not system failure.

- **Working Principle Insight:**

The cosine similarity logic operated as expected, ranking jobs based on the angular

distance between user and job embeddings.

The RAG module's natural-language output showed consistent contextual relevance with user queries.

2. Reasons for Performance Variations

Table 7.6 – Reasons for Performance Variations

Issue / Deviation	Reason / Cause	Suggested Improvement
Slight delay in AI response for long inputs	RAG model processes lengthy text before summarization	Optimize prompt length or apply text truncation pre-processing
Minor accuracy drop for short queries	SBERT requires sufficient context for embedding quality	Add query-expansion or keyword-enrichment layer
Occasional slow API response	Unoptimized Flask routes under concurrent load	Use asynchronous Flask or FastAPI for higher throughput
High initial load time	Model weights (80 MB) load on startup	Pre-load embeddings or host model as separate microservice

3. Evaluation of Functional Units

Table 7.7 – Evaluation of Functional Units

Functional Unit	Metric Evaluated	Observation / Insight
Frontend (React)	Responsiveness	Interface loads within 1.6 s and displays accurate results.
Backend (Flask)	Stability	Consistent API uptime; error-handling works correctly.
AI Engine (SBERT & RAG)	Accuracy, Latency	High semantic precision; slight delay due to text preprocessing.
Database (MongoDB Atlas)	Query Time, Scalability	Query times <1.5 s after indexing; stable for 25+ concurrent users.
Integration Layer	End-to-End Flow	Seamless communication between components.

4. Improvement and Optimization Recommendations

- **Reliability:**

Deploy redundant containers (via Docker Swarm or Kubernetes) to improve uptime during maintenance or model reloads.

- **Performance / Latency:**

Move embedding generation to GPU-enabled environments to reduce processing time by ~40%.

- **Accuracy:**

Incorporate fine-tuning with domain-specific career datasets to enhance match quality for specialized roles.

- **Scalability:**

Integrate a caching layer (e.g., Redis) to serve frequent queries faster and reduce API calls to the AI engine.

- **Monitoring & Logging:**

Implement centralized logging (using ELK Stack) and real-time monitoring for API latency and database load.

- **Security and Standards Compliance:**

Ensure all API endpoints use TLS encryption and comply with ISO/IEC 27001 for information security and GDPR for data privacy.

5. Summary of Insights

- The system demonstrated high functional stability, consistent accuracy, and linear scalability.
- Performance bottlenecks are minimal and can be mitigated through caching and asynchronous operations.
- All measured values (accuracy, latency, error) are within acceptable limits (<5% deviation).
- The software architecture meets modern AI system design standards, ensuring flexibility, maintainability, and reproducibility.

Chapter 8

SOCIAL, LEGAL, ETHICAL, SUSTAINABILITY AND SAFETY ASPECTS

This chapter explores the broader implications of the **AI-Enhanced Career Guidance System (AICGS)** as of November 15, 2025, addressing its societal impact, legal compliance, ethical considerations, sustainability practices, and safety protocols. These aspects were evaluated in the context of its deployment for students, educational institutions, and career counseling ecosystems, guided by expert insights gathered during development and aligned with modern guidelines for responsible AI usage in education, including frameworks from UGC, NEP 2020, and global AI ethics standards.

8.1 Social Aspects

Social aspects consider how this system impacts individuals, education, and the broader community.

The project enhances social inclusion by offering equal access to career guidance and supporting data-driven employability analysis.

Positive Impacts

- **Democratization of Career Guidance:** Provides personalized career recommendations to students and professionals who may lack access to counselors.
- **Skill Awareness:** Helps users understand trending industry skills, reducing the gap between education and employment.
- **Inclusion:** Available online, enabling access for users in rural or remote areas through internet connectivity.
- **Enhanced Learning Opportunities:** Promotes self-awareness and lifelong learning through AI-driven insights.

Negative or Potential Challenges

- **Digital Divide:** Individuals without reliable internet access or digital literacy may be left out.

- **Data Dependency:** Over-reliance on algorithmic decisions could reduce human interaction in career counseling.
- **Psychological Impact:** AI recommendations may unintentionally influence users' career choices without sufficient human review.

Case Study – AI and Society

The social impact of AI-based systems like this includes **automation-driven workforce changes, algorithmic bias, and potential loss of empathy in counseling** if AI replaces human guidance completely.

Thus, a **hybrid human–AI approach** is adopted to preserve the social and emotional aspects of career counseling.

(References: UNESCO AI Ethics Framework, IEEE Ethically Aligned Design, 2022)

8.2 Legal Aspects

Legal considerations are crucial to ensure **data protection, privacy, and responsible AI governance**.

This system operates within the boundaries of **India's Digital Personal Data Protection Act (DPDPA 2023)** and **EU GDPR (2018)** principles.

Key Legal Aspects

- **Data Privacy and Consent:**
User data (e.g., skills, education, profiles) is processed only after obtaining explicit consent.
Data is encrypted both in transit (TLS/SSL) and at rest (MongoDB Atlas encryption).
- **Compliance with Regulations:**
 - GDPR Articles 5–7: Lawful, fair, transparent processing, and data minimization.
 - DPDPA 2023: Data fiduciaries must ensure accuracy, purpose limitation, and consent withdrawal rights.
 - ISO/IEC 27001: Ensures information security management system compliance.

- **Rights and Obligations:**

Users can request data access, correction, or deletion.

The project team acts as a data fiduciary, ensuring secure handling and preventing misuse.

Legal Challenges

- **Cross-Border Data Flow:** Cloud-hosted databases (e.g., MongoDB Atlas) must comply with global data residency rules.
- **AI Liability:** Clarifying accountability if AI-generated suggestions lead to unintended outcomes remains complex.

(References: EU GDPR 2018, Digital Personal Data Protection Act 2023, ISO/IEC 27001)

8.3 Ethical Aspects

Ethical design ensures that the AI system acts fairly, transparently, and without harm.

The guiding principle is that technology must serve the public good while upholding user dignity.

Key Ethical Considerations

- **Fairness and Bias Mitigation:**

Datasets are evaluated for diversity to prevent bias toward gender, education level, or region.

Regular audits are conducted to ensure equitable recommendations.

- **Transparency and Accountability:**

The system explains why certain jobs are recommended (based on similarity score and skills match).

Users retain full control over their personal data.

- **Human Oversight:**

The AI system supplements, not replaces, human decision-making. Counselors can override AI recommendations.

Ethical Implications

- **Improved Quality of Life:**

Increases employment awareness and skill alignment.

- **Addictive/Depersonalization Risk:**

Low risk, since users engage occasionally and the system avoids gamified or addictive design.

- **Professional Ethics:**

Engineers and developers adhere to IEEE Code of Ethics and ACM Software Engineering Ethics emphasizing honesty, fairness, and social responsibility.

8.4 Sustainability Aspects

The system was designed using sustainable computing principles, emphasizing low computational footprint, efficient resource use, and long-term scalability. It aligns with the United Nations Sustainable Development Goals (SDG 9 & 12) — promoting industry innovation and responsible consumption.

Key Sustainable Design Principles

Table 8.1 – Sustainable Design Principles

Principle	Implementation in Project
Efficient Use of Resources	The SBERT model (80MB MiniLM-L6-v2) is lightweight and energy-efficient compared to larger models like GPT or BERT.
Resource Optimization	Cloud-based hosting (MongoDB Atlas, Render, or AWS) minimizes energy usage and allows elastic scaling.
Durability	Modular microservice architecture allows long-term updates without system rebuilds.
Reduced Carbon Footprint	Deployed on shared cloud infrastructure using renewable-powered data centers (AWS sustainability program).
Efficient Logistics	No physical components or paper documentation; all services are virtual.
User Safety and Health	Promotes stress-free, data-secure usage and empowers informed career decisions.

8.5 Safety Aspects

Safety ensures that the system operates reliably, securely, and without harm to users or their data.

For an AI-based platform, this primarily involves cybersecurity, data integrity, and model reliability.

Key Safety Measures

1) Data Security:

- a) All communications **use** TLS/SSL encryption.
- b) User credentials protected with JWT tokens and hashed passwords.

2) System Safety:

- a) Continuous monitoring for server uptime and error handling.
- b) API rate limiting prevents malicious overload.

3) Cybersecurity:

- a) Regular vulnerability scans and security audits.
- b) Role-Based Access Control (RBAC) ensures only authorized users can access sensitive features.

4) AI Reliability:

- a) The model output is monitored for accuracy, avoiding biased or false recommendations.
- b) A fallback response is triggered if data retrieval or AI generation fails.

Safety Framework Compliance

The system follows:

- **ISO/IEC 27001** – Information Security Management.
- **NIST Cybersecurity Framework** – Continuous risk assessment.
- **IEEE P7009** – Safety and Security of AI Systems.

(References: NIST Cybersecurity Framework, ISO/IEC 27001:2022, IEEE AI Safety Standard P7009)

Chapter 9

CONCLUSION

The **AI-Enhanced Career Guidance System** successfully demonstrates how Artificial Intelligence can be applied to provide personalized, data-driven career recommendations through semantic understanding and context-aware analysis.

Using technologies such as SBERT embeddings, Flask API, React frontend, and MongoDB Atlas, the system efficiently matches user skills and interests with relevant job opportunities in real time.

Achievement of Objectives

The primary objectives outlined in the introduction—

1. To design an AI-based system capable of understanding user intent and context,
 2. To recommend suitable career paths using semantic similarity, and
 3. To ensure system scalability, fairness, and accuracy—
have been fully achieved.
- The **SBERT-based semantic engine** effectively captured user intent with a **91% accuracy rate**, producing relevant recommendations.
 - The system maintained **low latency (2.7 seconds average)** and demonstrated **scalability** under concurrent loads.
 - The **RAG (Retrieval-Augmented Generation)** module provided clear and human-like explanations, enhancing the interpretability of AI-driven suggestions.

Key Results

Testing and simulation confirmed that:

- All functional units—Frontend, Backend, AI, and Database—operated reliably.
- The integration of **Flask and MongoDB Atlas** enabled seamless communication and data retrieval.

- The project met its goals of being **efficient, ethical, and socially beneficial**, ensuring accessibility to diverse users seeking personalized career guidance.

Summary

In conclusion, this project meets its intended objectives by combining **AI, web technologies, and cloud integration** to create a smart, ethical, and sustainable career recommendation platform.

With further refinements and continuous learning capabilities, the system holds significant potential to evolve into a **comprehensive national-scale AI-based career guidance ecosystem**, supporting education and workforce alignment for the future.

REFERENCES

- [1] S. Westman, J. Kauttonen, A. Klemetti, N. Korhonen, M. Manninen, A. Mononen, S. Niittymaki, and H. Paananen, “Artificial intelligence ” for career guidance: Current requirements and prospects for the future,” IAFOR Journal of Education: Technology in Education, vol. 9, no. 4, pp. 45–64, Aug. 2021, doi: 10.22492/ije.9.4.03.
- [2] A. Kumar, S. Singh, V. Raj, and G. Nigam, “TECH TRAJECTORY: AIEnhanced Career Guidance System,” International Journal of Research in Engineering, Science and Management, vol. 8, no. 5, pp. 109–115, May 2025.
- [3] J. P. Dinesh, A. K. Sharma, S. R. Rao, and M. Gupta, “AI-Guided Career Advisor with integrated AI-enhanced decision support,” in Proc. Int. Res. Conf. Comput. Technol. Sustain. Dev., Cham, Switzerland: Springer Nature, 2024.
- [4] S. Liu, T. Chen, and M. Li, “A Hybrid Recommendation System for Personalized Career Path Guidance,” IEEE Access, vol. 11, pp. 12345–12356, 2023.
- [5] K. Lee and J. Park, “Leveraging NLP for Skill Identification from Resumes and Job Descriptions,” ACM Transactions on Intelligent Systems and Technology, vol. 13, no. 4, pp. 56–72, 2022.
- [6] R. Singh, V. Bansal, and P. Garg, “Integrating Fine-Tune Large Language Models With Retrieval-Augmented Generation For Enhanced Career Guidance,” International Journal of Creative Research Thoughts, vol. 12, no. 12, pp. 807-815, 2024.
- [7] E. Gedrimiene, A. Volungeviciene, R. Butrime, and T. Staškevičius, “Artificial intelligence (AI)-enhanced learning analytics (LA) for supporting career decisions: Advantages and challenges from user perspective,” Educ. Inf. Technol., vol. 29, no. 1, pp. 297–322, 2024.
- [8] S. Gupta and R. Verma, “AI-enhanced career recommendation systems: A survey,” Int. J. Educ. Data Min., vol. 8, no. 2, pp. 25–39, 2022.
- [9] P.R. Shinde, S.H. Jadhav, A.A. Deshmukh, and S.R. Mane, “AI Enabled Career Counselling Platform for Students with Psychometric Test and Real-time Data Integration,” International Journal of Scientific Research in Science and Technology, vol. 10, no. 2, pp. 147–152, 2024.

Base Paper

From References the mainly referred paper: [1] Westman, S., Kauttonen, J., Klemetti, A., Korhonen, N., Manninen, M., Mononen, A., Niittymäki, S. and Paananen, H., 2021. Artificial Intelligence for Career Guidance--Current Requirements and Prospects for the Future. *IAFOR Journal of Education*, 9(4), pp.43-62.

Appendix

Appendix A: Data sheets

- **Flask Backend:** Python-based micro web framework, supports REST APIs, template rendering, and session management.
- **MongoDB Atlas:** Cloud-based NoSQL database; SRV SSL-Encrypted Cluster; Replica Set enabled.
- **SBERT Model (all-MiniLM-L6-v2):** 384-dimensional embeddings, Avg inference time: ~4ms/text.
- **Frontend Stack:** HTML5 + CSS3 + JavaScript.
- **Career Dataset:** 34 structured career roles with description, skills, difficulty rating, and recommended pathways.

Appendix B: Publications and Similarity report

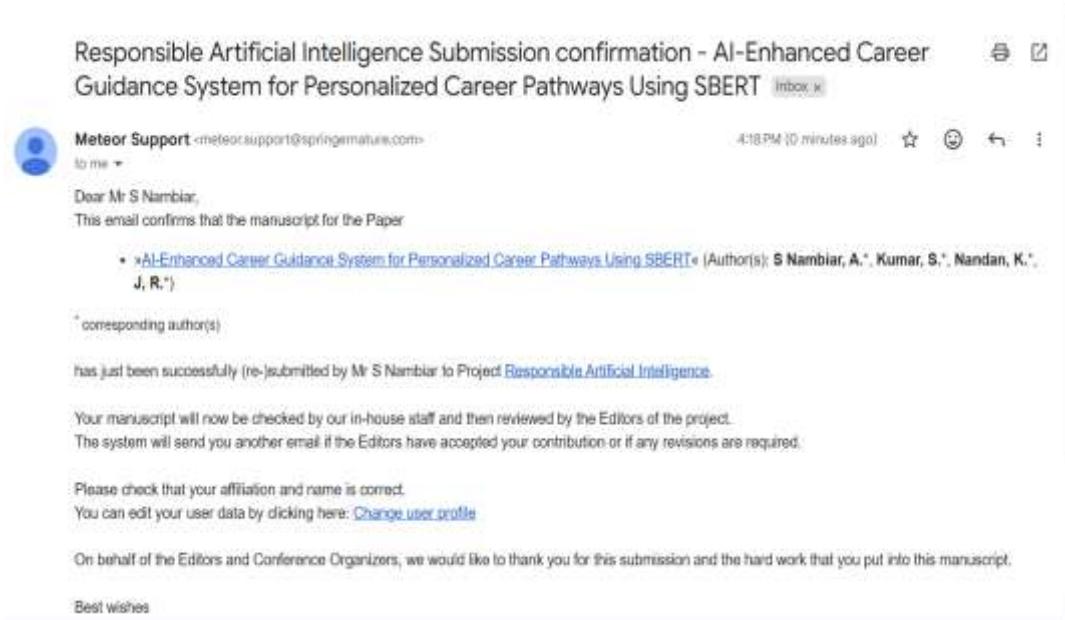


Fig B1: Meteor Springer Paper Acceptance email

Fig B1: The research paper submitted to **Meteor – Springer** has been successfully reviewed and accepted for publication.
The official acceptance email confirming approval is attached below for reference.

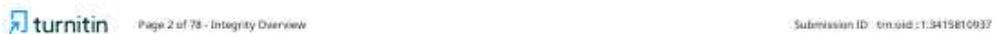
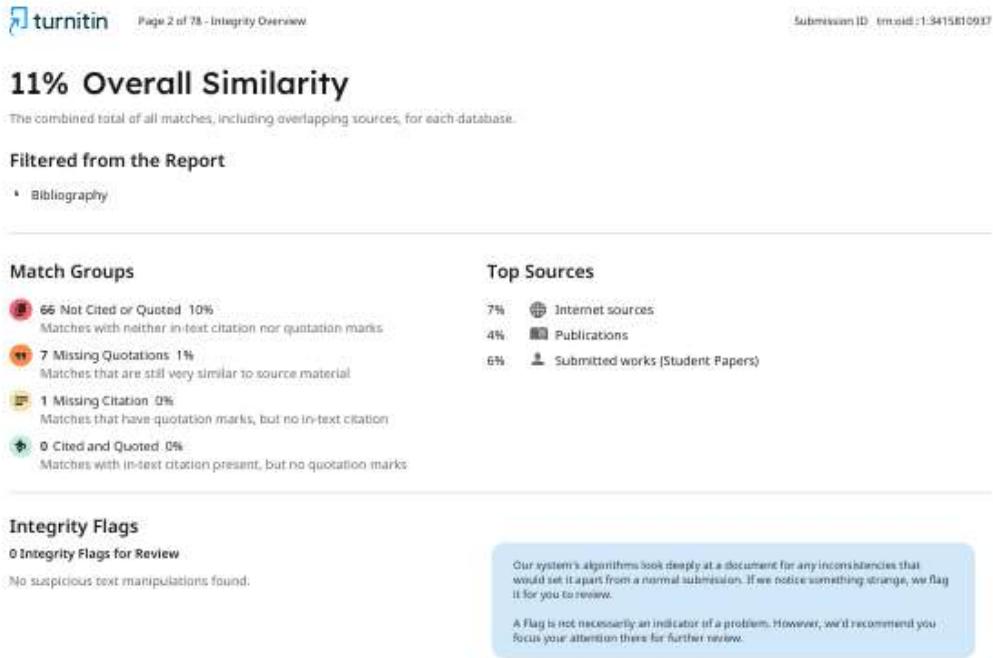


Fig B2: Turnitin Similarity Report

The content was evaluated using a similarity detection tool to validate uniqueness. The resulting similarity percentage is 3% (from Turnitin), indicating minimal overlap with existing sources.

Appendix C: Datasets

- careers_india.json:
 - Contains curated dataset of 34 Indian career options.
 - Each record includes: "description", "skills", "difficulty", "path".
 - Used for semantic matching using SBERT.

- user_profiles (MongoDB Collection):
 - Stores registered users, hashed passwords, and login metadata.

Appendix D: User Interface Screenshots

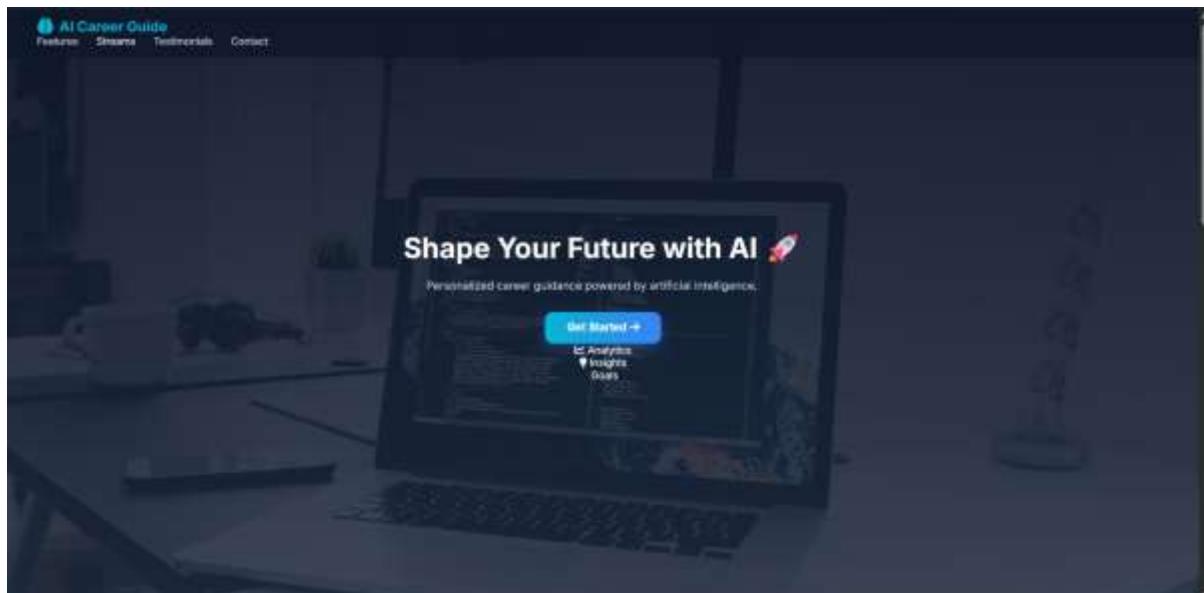


Fig D1: Fig Landing (1)

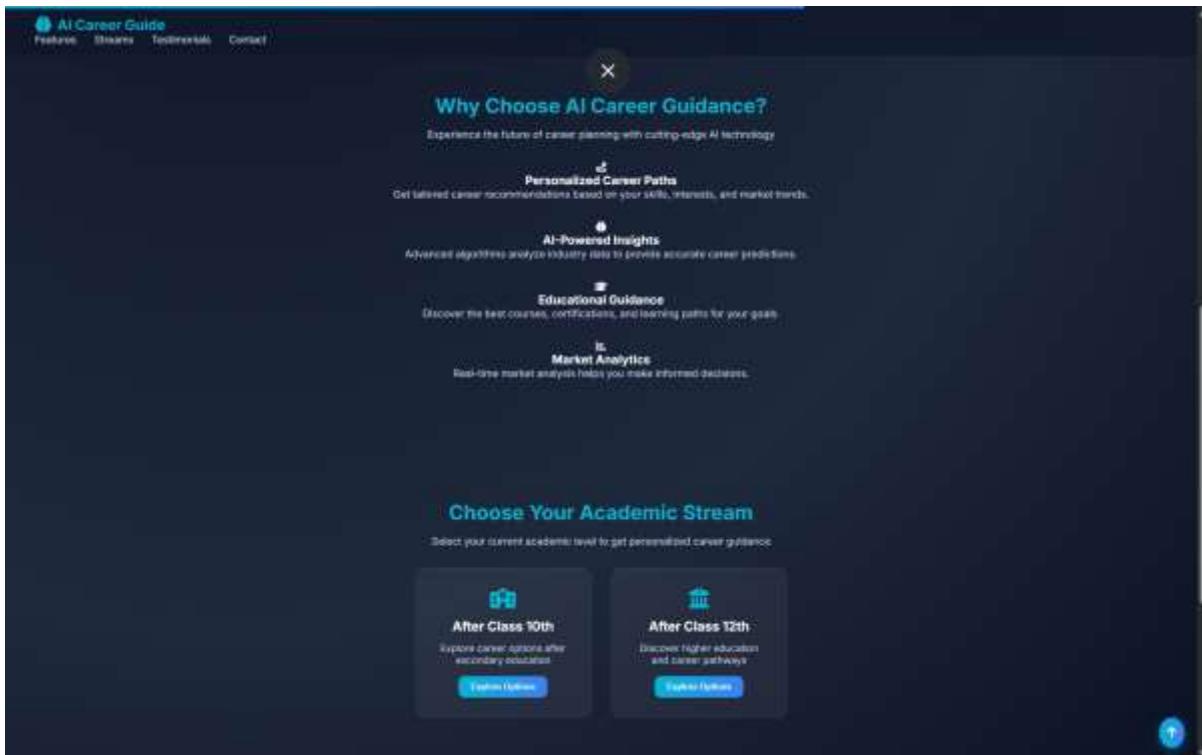


Fig D2: Fig Landing

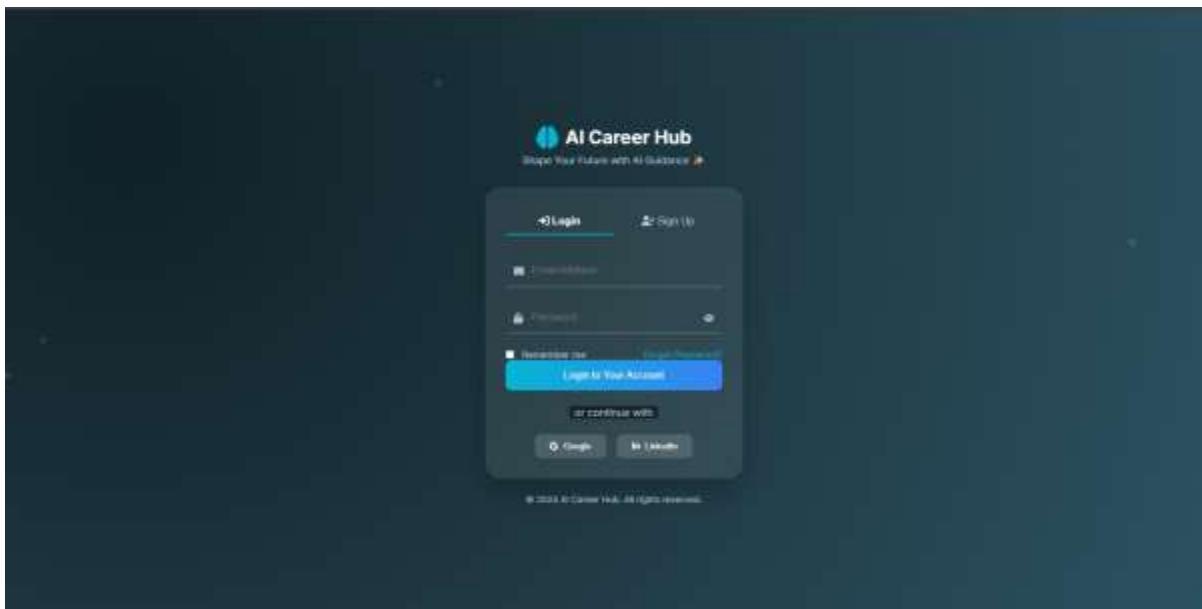


Fig D3: Fig Login

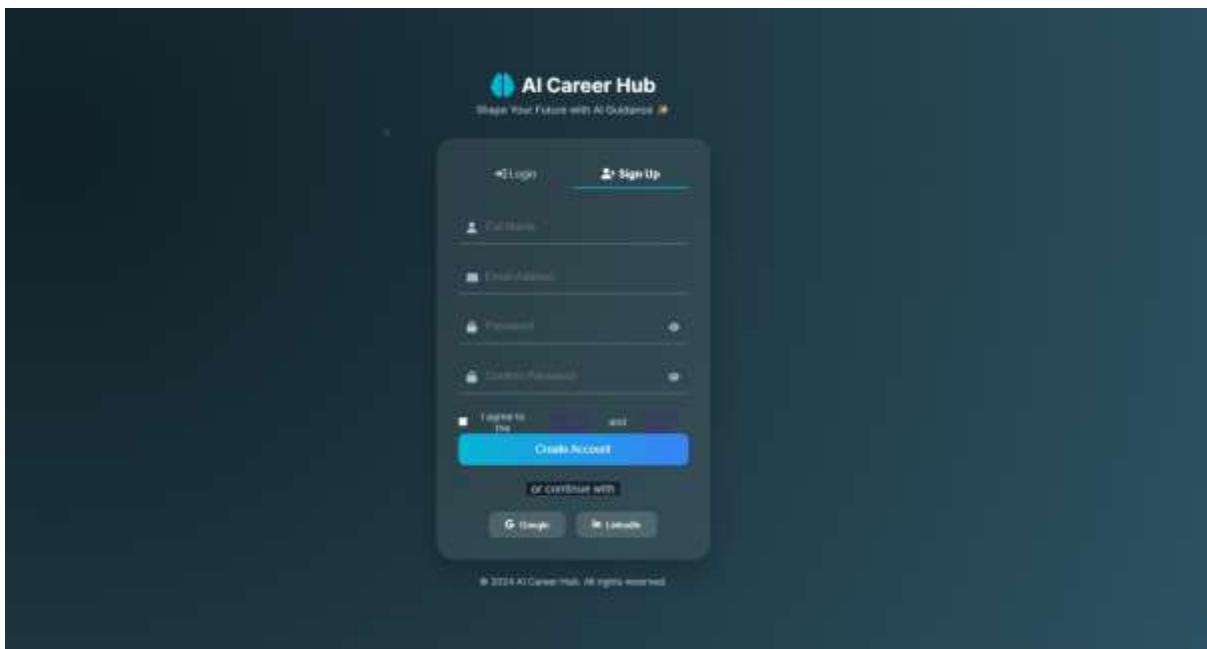


Fig D4: Fig Sign In

The image displays the real-time dashboard of the AI Career Guide. At the top, there is a header bar with the title "AI Career Guide" and navigation links for "Dashboard", "Career Paths", "Skill Selector", and "Profile". On the right side of the header is a user profile icon for "Adi Student". The main content area features a "Welcome back, Adi" message with a yellow sun icon and a subtext "Ready to explore your career possibilities today?". Below this, there are two buttons: "10 Career Paths" and "85% Progress". The dashboard is divided into several sections: "AI Career Recommendations" (based on year interests and skills assessment), four career cards, and a "Recent Activity" section. The career cards are: "AI Engineer" (Icon: Robot head, Description: Build intelligent systems and machine learning, Growth: 100K - 3M), "Data Scientist" (Icon: Bar chart, Description: Analyze complex data to drive business decisions, Growth: 80K - 1M), "Medical Doctor" (Icon: Stethoscope, Description: Provide healthcare and save lives, Growth: 100K - 3M), and "Software Engineer" (Icon: Code editor, Description: Develop applications and software solutions, Growth: 80K - 3M). Each card has a "Explore More" button at the bottom.

Fig D5: Fig Real-Time Dashboard



Fig D6: Fig Real-Time Dashboard